Fisher-Glass Learning: Classification Bounds and Subpopulation Audits

El Mustapha Mansouri
School of Engineering, Institute of Science Tokyo
Yokohama, Kanagawa 226-8501, Japan
mansouri.e.2224@m.isct.ac.jp

Abstract

High predictive accuracy can coexist with sensitivity to task-preserving nuisance changes. Building on the Fisher-glass framework, we derive local classification lower bounds and develop subpopulation audits based on nuisance-projected information. The bounds connect efficient task information to minimax excess classification risk under explicit distinguishability and decision-separation conditions. We derive regularization-aware audit bounds and evaluate frozen population gates on held-out data. A prospective Waterbirds study uses an uncentered feature–Schur diagnostic constructed from task logits and background-probe scores to select an eligible validation group. Across five training seeds, interventions preserving the extracted bird foreground and labels yield 93.2% accuracy under same-type controls and 40.6% under opposite-type swaps, a 52.6-percentage-point difference. Search using task labels and a fixed nuisance direction also identifies held-out populations with high accuracy and low estimated projected predictive Fisher information. Training and acquisition experiments evaluate changes in information and predictive performance together. The resulting framework pairs predictive performance with explicit measures of task information and nuisance alignment.

1 Introduction

Reliable prediction under distribution shift requires understanding the evidence supporting a model's decisions. A classifier can perform well on a subpopulation by using a nuisance cue correlated with its label, then fail when that correlation changes. Group-robust learning and error-based reweighting make subpopulation performance central to robust machine learning [Sagawa et al., 2020, Liu et al., 2021]. We ask a complementary question: *how much task information remains after nuisance variation is accounted for?*

We investigate this question through a prospective Waterbirds intervention. Using validation data, we select an eligible group by the uncentered feature–Schur fraction of task-logit and background-probe coordinates. Models and the selected group are frozen before held-out evaluation. Across five training seeds, the selected group achieves 92.0% accuracy on untouched images. Background replacements preserving the extracted bird foreground and labels yield 93.2% accuracy under same-type controls and 40.6% under opposite-type swaps, a controlled difference of 52.6 percentage points. This comparison quantifies sensitivity to the audited background change within an accurately classified population.

Efficient-score theory describes the task information remaining after nuisance adjustment through an orthogonal projection [Bickel et al., 1993, Tsiatis, 2006]. The Fisher-glass framework of Mansouri and Arai [2026] connects this local geometry to inverse-information tails and a population tail certificate. We build on this framework to study classification difficulty and the evaluation of trained

40th Conference on Neural Information Processing Systems (NeurIPS 2026).

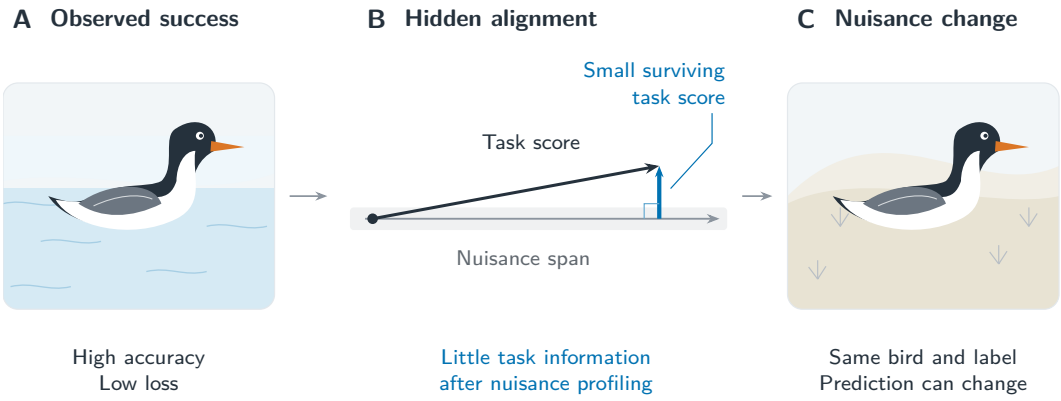


Figure 1: **Conceptual motivation for auditing task information under nuisance variation.** A nuisance cue can support high accuracy and low loss (A), while task–nuisance alignment leaves little information after projection (B). A task-preserving background change can expose the resulting sensitivity (C). The identical vector bird is reused in both environments; these are conceptual illustrations, not experimental images.

predictors. Figure 1 introduces the geometric motivation, and Section 2 specifies the task and nuisance coordinates.

Our classification result combines local indistinguishability with quantitative Bayes-decision separation. The resulting lower bound retains the environment-dependent information, separation constants, and local radius. We then develop the statistical ingredients needed for an operational audit: regularization-aware information bounds, supported population gates, and independent evaluation of selected subpopulations. Given task labels and a fixed nuisance direction, the search identifies a held-out Waterbirds population with 98.8% accuracy and low estimated projected predictive Fisher information.

Our contributions are:

- **Local classification theory.** We derive an excess-classification-risk lower bound under explicit distinguishability and decision-separation conditions, connecting efficient task information to local learning difficulty.
- **Statistical audit methodology.** We derive bounds for regularized Schur estimates and independently evaluated population gates, accounting for membership normalization and support correction.
- **Machine-learning evaluation.** We study controlled background sensitivity, held-out discovery of accurate low-information populations, and the joint information–performance outcomes of targeted training and acquisition.

2 Task–Nuisance Geometry and Information Collapse

Let $e \sim \Pi$ index an observed group or a latent subpopulation. We fix the task contrast, admissible nuisance variations, and a common normalization of task coordinates throughout the audit. Information then describes the task distinction remaining after adjustment for those nuisance variations.

2.1 Task Information After Nuisance Profiling

Consider a local model $P_{\theta,\mu,e}$ with task parameter θ and nuisance parameter μ. A likelihood score measures the local change in log likelihood as a parameter varies. At $P_e = P_{0,0,e}$, assume differentiability in quadratic mean, with centered, square-integrable task-score vector $s_{T,e}$ and nuisance-score vector $s_{N,e}$. Let $\mathcal{T}_{N,e}$ be the closed span of admissible nuisance scores in $L^2(P_e)$, with inner product $\langle f, g\rangle_e = \mathbb{E}_{P_e}[fg]$, and let $\Pi_{\mathcal{T}_{N,e}}$ denote orthogonal projection onto that span.

The **efficient task score** and its information are

$$\begin{aligned} s^{\mathrm{eff}}_{T,e} &= s_{T,e} - \Pi_{\mathcal{T}_{N,e}} s_{T,e}, \\ F_e := F^{\mathrm{eff}}_{T|N,e} &= \mathbb{E}_{P_e}[s^{\mathrm{eff}}_{T,e}(s^{\mathrm{eff}}_{T,e})^\top] = J_{TT,e} - J_{TN,e} J^{-1}_{NN,e} J_{NT,e}. \end{aligned} \tag{1}$$

The projection acts componentwise. The Schur expression applies when finite nuisance scores span $\mathcal{T}_{N,e}$ and $J_{NN,e} \succ 0$, with $J_{AB,e} = \mathbb{E}_{P_e}[s_{A,e} s_{B,e}^\top]$ [Bickel et al., 1993, Tsiatis, 2006, Cox and Reid, 1987]. Write $\lambda_e = \lambda_{\min}(F_e)$ for the weakest task information. A zero eigenvalue identifies a task direction contained in the nuisance span.

For scalar task and nuisance scores with positive raw information, their Fisher correlation $\rho_e = J_{TN,e}/\sqrt{J_{TT,e} J_{NN,e}}$ gives

$$F_e = J_{TT,e}(1 - \rho_e^2), \qquad \underbrace{F_e / J_{TT,e}}_{\text{surviving-information fraction}} = 1 - \rho_e^2. \tag{2}$$

Absolute information measures remaining curvature in the declared coordinates. The fraction measures the proportion retained, lies in $[0, 1]$, and is invariant to nonzero scalar rescaling of either score. It compares alignment across slices with different raw information; its feature–Schur analogue supplies the prospective selection statistic. The model-level Fisher-Glass certificate, defined below, instead summarizes absolute information across environments.

2.2 Task–Nuisance Tangency and Loss–Information Separation

Let b_e, r_e be centered, orthogonal, unit-norm scores, with nuisance span $\mathrm{span}\{b_e\}$. For $\delta_e \geq 0$, the normalized construction gives

$$a_e = b_e + \delta_e r_e, \quad J_{TT,e} = 1 + \delta_e^2, \quad F_e = \delta_e^2, \quad \frac{F_e}{J_{TT,e}} = \frac{\delta_e^2}{1 + \delta_e^2}. \tag{3}$$

In the transverse-information geometry of Mansouri and Arai [2026], projection removes b_e, leaving $\delta_e r_e$. Raw information stays above one as efficient information vanishes quadratically; δ_e measures distance from the nuisance span. When the nuisance feature also supports small observed prediction risk, low loss coexists with arbitrarily weak efficient information. This separation concerns the evidence supporting prediction, even within accurately classified populations. The formal construction and proof appear in Appendix A.1. Figure 2 illustrates the geometry's population consequence: frequent near-tangencies produce a heavy inverse-information tail.

3 Tail Identifiability and Local Learning Difficulty

The frequency of near-null environments determines the inverse-information tail. Its classification consequence additionally depends on how weak task alternatives change the Bayes decision. We separate these ingredients before defining the certificate.

3.1 Collapse Geometry Determines Information Tails

In the tangency construction, $\lambda_e = \delta_e^2$: if $\Pr(\delta_e < t) \sim Ct^\gamma$, then $\Pr(\lambda_e^{-1} > x) \sim Cx^{-\gamma/2}$. The geometric tail law of Mansouri and Arai [2026] relates the inverse-information exponent to null-set codimension, information vanishing order, and nearby probability mass. We state the law with explicit regularity and localization assumptions.

Theorem 1 (Geometric tail law [Mansouri and Arai, 2026])**.** *Let $z \in \mathbb{R}^d$ parameterize environments with density p, and let the Fisher-null set $\mathcal{Z} = \{z : \lambda(z) = 0\}$ be a nonempty compact C^2 embedded submanifold without boundary, of codimension $c \in \{1, \ldots, d\}$. Suppose λ is continuous on the closure of a fixed tubular neighborhood of $\mathcal{Z}$ and bounded away from zero outside it. In normal coordinates $z = s + r\omega$, let $s \in \mathcal{Z}$, r be normal distance, and ω a unit normal direction. As $r \downarrow 0$, assume uniformly that information vanishes as $a(s,\omega) r^{2m}(1 + o(1))$ and density as $b(s,\omega) r^\nu (1 + o(1))$, where $m > 0$, $\nu > -c$, and the continuous coefficients are uniformly positive and bounded. Then, for some $C > 0$,*

$$\Pr(\lambda(z)^{-1} > x) \sim Cx^{-\beta}, \qquad \beta = \frac{c + \nu}{2m}. \tag{4}$$

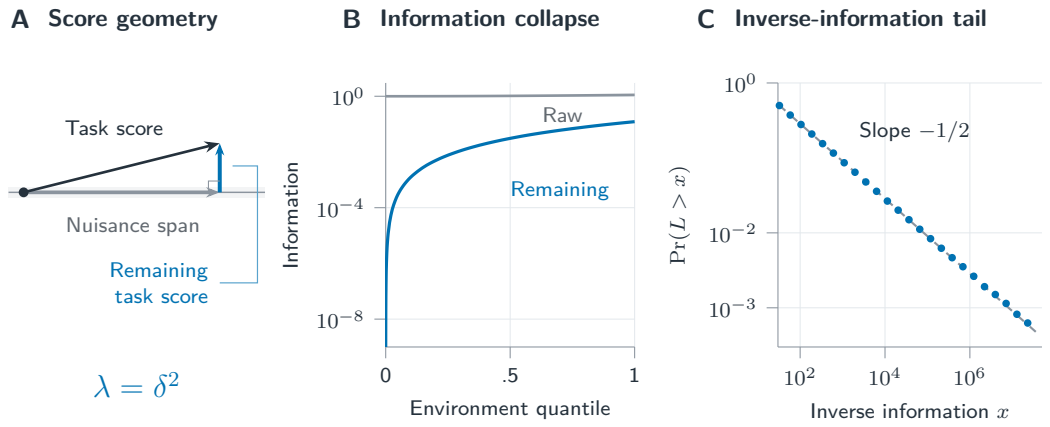


Figure 2: **Synthetic illustration of the geometric tail law** [Mansouri and Arai, 2026]. (A) The surviving score has norm δ and information $\lambda = \delta^2$. (B) Raw information stays above one while projected information approaches zero. (C) Empirical inverse-information survival (blue) follows the theoretical $x^{-1/2}$ tail (dashed). B–C share one regenerated seed-11 draw of 300,000 environments, $\delta \sim \mathrm{Uniform}(0, 0.35)$ and $L = (\delta^2 + 10^{-12})^{-1}$. The separately reported five-seed exponent is 0.498 ± 0.003 (nominal 95% normal-approximation interval, ± 1.96 standard errors).

An information threshold defines a narrow tube around $\mathcal{Z}$; integrating its probability mass determines the exponent. Greater codimension makes simultaneous collapse rarer, whereas higher-order vanishing produces a heavier tail. Quadratic collapse with positive density gives $\beta = c/2$, and the exact inverse-information moment is finite precisely when $\beta > 1$. Thus substantial information across most environments can coexist with divergent average inverse information. Appendix A.2 gives a complete proof under these assumptions, including the asymptotic constant.

3.2 Efficient Information and Local Classification Difficulty

The geometric tail acquires a classification consequence when nearby nuisance-adjusted alternatives induce quantitatively different optimal decisions. Define classification risk $R_P(g) = P\{g(X) \neq Y\}$, Bayes risk $R_P^\star = \inf_g R_P(g)$, and excess risk $\mathcal{E}_P(g) = R_P(g) - R_P^\star$. For a specified local family, let $\mathfrak{R}_e(n)$ denote minimax expected excess classification risk based on n independent observations from environment e.

Theorem 2 (Local classification difficulty)**.** *Fix $\lambda_e > 0$ and a unit weakest efficient-information direction u_e. For every $0 < t \leq t_{0,e}$, suppose the local family contains $P_{e,+t}, P_{e,-t}$ with task coordinates $\pm t u_e$ and nuisance coordinates allowed to vary. Assume, for constants $K_e, c_e, t_{0,e} > 0$ and $q > 0$,*

$$\mathrm{KL}(P_{e,+t} \| P_{e,-t}) \leq K_e t^2 \lambda_e, \qquad \inf_g \{\mathcal{E}_{P_{e,+t}}(g) + \mathcal{E}_{P_{e,-t}}(g)\} \geq c_e t^q. \tag{5}$$

Then, for a universal constant $c_0 > 0$,

$$\mathfrak{R}_e(n) \geq c_0 c_e \min\{t_{0,e}^q, (n K_e \lambda_e)^{-q/2}\}. \tag{6}$$

The KL condition limits distinguishability after nuisance adjustment; decision separation specifies the cost of confusing the alternatives. Both must hold over the environment-specific radius $t_{0,e}$. For quadratic separation, $q = 2$, the interior bound scales as $c_e/(n K_e \lambda_e)$. The minimum retains the finite scale of the local problem as information decreases. This is a conditional local result: efficient

information sets the testing scale, while boundary geometry supplies its classification meaning. The complete testing proof and its tail-aggregation implication are in Appendix A.3.

3.3 The Fisher-Glass Certificate

Upper-tail conditional value-at-risk, CVaR_α, averages the largest $1-\alpha$ fraction of a quantity, with fractional weighting at atoms [Rockafellar and Uryasev, 2000, 2002]. Following Mansouri and Arai [2026], we apply the Fisher-glass certificate to classification environments $e \sim \Pi$, using the weakest task information λ_e and stabilization ε:

$$L_e^{(\varepsilon)} = (\lambda_e + \varepsilon)^{-1}, \qquad I_\alpha^{\mathrm{FG}} = \left[\mathrm{CVaR}_{\alpha, e\sim\Pi}(L_e^{(\varepsilon)})\right]^{-1}, \quad 0 < \alpha < 1. \tag{7}$$

A larger certificate means a smaller inverse-information tail, expressed in information units. Stabilization $\varepsilon > 0$ caps inverse loss at $1/\varepsilon$; it does not remove the underlying null geometry. Under the theorem's power-law regime, the exact limit $\varepsilon \downarrow 0$ is positive for $\beta > 1$ and vanishes for $\beta \le 1$.

Proposition 1 (Information coverage)**.** *For $\rho > \varepsilon$,*

$$I_\alpha^{\mathrm{FG}} \ge \rho \quad \Longrightarrow \quad \Pr_{e\sim\Pi}(\lambda_e \ge \rho - \varepsilon) \ge \alpha. \tag{8}$$

The certificate therefore establishes an absolute information floor on at least an α-fraction of the specified environment distribution. Its connection to classification retains the local-radius and decision-separation factors in the preceding bound. The CVaR representation, coverage proof, and limiting interpretation appear in Appendix A.4. The operational audit estimates this information target on independently evaluated populations.

4 Computing and Auditing Fisher-Glass

An operational audit fixes the task contrast, nuisance hypothesis, matrix construction, and audit distribution before measuring information on independent data. We work in low-dimensional task and nuisance subspaces, so the reported information concerns these declared coordinates. Coordinate normalization is specified for each experiment; fitted nuisance probes and temperature calibration are frozen before final evaluation. The audit reports support and conditioning alongside its information values.

4.1 Operational Scores and Nuisance Coordinates

The audit uses a declared matrix target. Predictive Fisher is computed from likelihood scores under the fitted predictive distribution; observed-label score moments and unweighted feature moments define additional audit targets. The prospective Waterbirds experiment uses the uncentered feature–Schur fraction of fixed task-logit and background-probe coordinates. Hidden-slice discovery uses estimated temperature-calibrated predictive Fisher with class-specific shrinkage $\eta = 0.10$, retained during search and held-out evaluation. Unweighted feature geometry supplies a complementary comparison. Appendix B.1 specifies each construction.

The experiments prespecify nuisance coordinates from metadata, fitted annotation probes, or a fixed semantic direction. Without nuisance annotations, the specified semantic direction serves as the audited nuisance hypothesis. We report fitting splits, normalization, and nuisance-block conditioning; cross-validation assesses gate stability. Appendix D.5 records the checks performed. Formal anchor task-preservation tests and admission thresholds are criteria for extending this family, rather than implemented selection steps.

4.2 Stable Estimation and Exact-Information Audits

Let M_i be the per-example task–nuisance matrix under the chosen construction, with task and nuisance blocks indexed by T and N. For a predefined slice g of size n_g, form $\widehat{J}_g = n_g^{-1} \sum_{i\in g} M_i$ and compute the ridge Schur estimate

$$\widehat{F}_g^{\mathrm{rid}} = \widehat{J}_{TT,g} - \widehat{J}_{TN,g}(\widehat{J}_{NN,g} + \lambda_{\mathrm{rid}} I)^{-1} \widehat{J}_{NT,g}. \tag{9}$$

Algorithm 1 Fisher-Glass audit

1: **Specify:** predictor, matrix, coordinates, normalization, calibration, and audit weights.
2: **Discover and freeze:** predefined slices or constrained gates fitted on discovery data.
3: **Audit independently:** check support and conditioning; compute ridge, unregularized, and lower values.
4: **Report:** raw and projected information, surviving fraction, slice support, tail certificate, and uncertainty.

Real-data computations use $\lambda_{\mathrm{rid}} = 10^{-3}$. Ridge stabilizes residualization but raises information relative to the unregularized target. Let a_g bound estimation error and d_g bound the ridge-to-unregularized gap on a simultaneous event over the audited slices. For declared weights Π_G, use

$$\begin{aligned} \underline{\lambda}_g &= [\lambda_{\min}(\widehat{F}_g^{\mathrm{rid}}) - a_g - d_g]_+, \\ \underline{I}_{\alpha,G} &= \left[\mathrm{CVaR}_{\alpha,g\sim\Pi_G}\left((\underline{\lambda}_g + \varepsilon)^{-1}\right)\right]^{-1}. \end{aligned} \tag{10}$$

Here $[x]_+ = \max\{x, 0\}$. On that event, this lower-bounds the unregularized population certificate for the declared matrix and audit distribution. Certification also requires a positive lower confidence bound on nuisance conditioning. The implemented scalar audits use Bonferroni-adjusted Gaussian/Wald moment simulation at nominal 95% confidence, conditional on the fitted model and evaluated slices. Appendix B records the gap identity, error allowances, and numerical comparison.

4.3 Hidden-Slice Discovery and Held-Out Certification

A gate $h_\phi(x, y) \in [0, 1]$ assigns membership weights. Discovery maximizes inverse projected information subject to effective-sample-size, maximum-weight, and class-mass constraints, with shrinkage toward the same-class pooled matrix. Five validation folds assess transfer and stability; after fixing the protocol, a final gate is fitted on all discovery data, frozen, and evaluated on an untouched split. Its unshrunk target is the weighted ratio $J_\phi = \mathbb{E}[h_\phi M]/\mathbb{E}[h_\phi]$, with positive membership mass. Uncertainty must account for numerator, normalization, and any constraint-induced weight adjustment; the reported Waterbirds target retains the stated shrinkage. Appendix C specifies the constraints, objective, and held-out target.

5 Evaluating Information Audits

We test the diagnostic on a prospective nuisance intervention, isolate its geometry in synthetic experiments, and discover low-information populations on held-out data. The corresponding quantities are the surviving fraction, inverse-information tail, and predictive projected Fisher.

5.1 Prospective Diagnosis Under a Nuisance Intervention

On Waterbirds [Sagawa et al., 2020], the four groups cross bird class with background type and are written *class/background*. We fully fine-tune ImageNet-initialized MobileNetV3-small models [Howard et al., 2019] with empirical risk minimization. Validation data fit the temperature and cross-fitted background probe. Eligible slices have accuracy at least the overall validation accuracy and negative log-likelihood at most the median across slices. Among them, we select the smallest surviving fraction, $\lambda_{\min}(\widehat{F}_g^{\mathrm{rid}})/\widehat{J}_{TT,g}$. After pilot experiments, this rule is fixed for five additional training seeds, 22–26. Each checkpoint, temperature, probe, selected slice, and audit constant is frozen before intervention scoring.

The rule selects **waterbird/water in all five seeds**: its validation surviving fraction is 24.8%, compared with 59.0% for landbird/land, the other eligible slice. Held-out interventions recompose birds onto same-type or opposite-type backgrounds using the same extracted foreground layer, matte, and task labels. Same-type recomposition controls for the image-construction procedure. Selected-slice accuracy is 92.0% on original images, 93.2% under the control, and 40.6% under opposite-type replacement: a 52.6 ± 1.4 **percentage-point controlled drop**, with a nominal 95% normal-approximation interval across seeds.

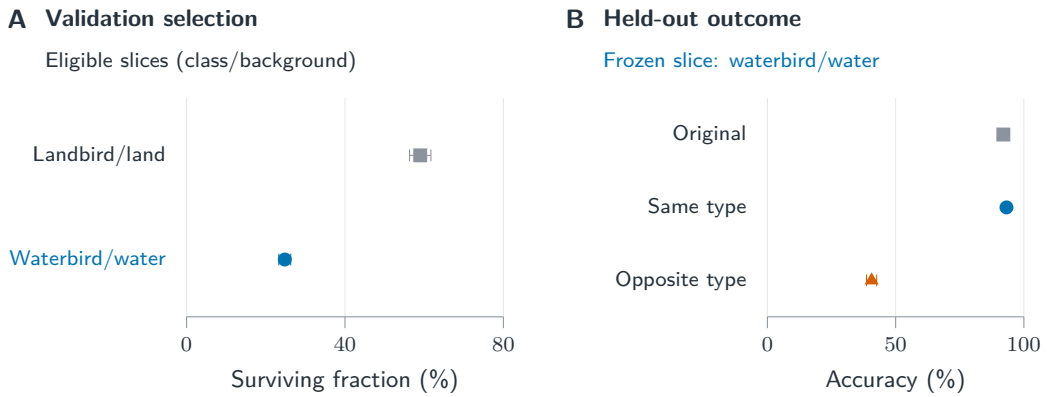


Figure 3: **Validation selection anticipates held-out background sensitivity.** (A) Waterbird/water has the lower uncentered feature–Schur fraction among eligible slices and is selected in all five seeds. Matched minimum-accuracy and maximum-NLL selectors choose the same group in all five. Eligibility and selection are evaluated separately within each seed; names denote class/background. (B) With this selection frozen, ERM accuracy is high on original images and same-type recompositions but falls under opposite-type background recomposition, a controlled difference of 52.6 percentage points. Intervals are nominal 95% normal approximations (± 1.96 standard errors) across seeds 22–26; the same-type arm controls for recomposition.

Within the eligible set, minimum surviving fraction, minimum validation accuracy, and maximum validation NLL select waterbird/water in all five seeds, with a 52.6-point controlled accuracy drop (Figure 3). Minimum raw and absolute projected information select landbird/land, with a 26.7-point drop. Centering switches selection in all five seeds, identifying slice means as a ranking contributor; predictive Fisher retains waterbird/water in all five. Appendix D.2 reports these checks, the complete selector comparison, and individual flip prediction.

5.2 Synthetic Classification and Tail Checks

Using nuisance feature s, orthogonal residual r, and task coordinate $z = s + \delta_e r$ as in Section 2, the evaluated classifier has logit $4s$. For near-null mass exponent $\gamma = 1$, five-seed mean accuracy is 94.54% and loss-$\mathrm{CVaR}_{0.95} = 0.262$. Yet $I^{\mathrm{FG}}_{0.95} = (1.22 \pm 0.53) \times 10^{-8}$ with $\varepsilon = 10^{-12}$, despite raw information bounded below by one.

The inverse-information exponent is 0.498 ± 0.003, matching $\beta = \gamma/2 = 0.5$ in the geometric tail law of Theorem 1 (Figure 2). Log-survival regression uses the largest 5% of inverse-information observations. Independent-branch and reserve experiments check the structural laws of Mansouri and Arai [2026], including $\beta = K/2$ for K branches. Uncertainties are nominal 95% normal-approximation intervals across five seeds. These measurements exhibit favorable prediction alongside weak projected information and verify the predicted inverse-information tail behavior. Appendix D.3 records the structural checks and additional summaries.

5.3 Discovering Low-Information Subpopulations

Discovery receives images, task labels, frozen ImageNet MobileNetV3-small representations, and a fixed CLIP water-versus-land nuisance direction [Radford et al., 2021]. Background and group annotations remain sealed through discovery and test evaluation. A class-conditional linear gate targets 10% within-class mass under the support constraints of Section 4.3. Five validation folds assess transfer; after fixing the protocol, the gate is refit on validation data, frozen, and evaluated once on the untouched test split.

The discovered slice has 98.8% **accuracy** and lower estimated projected predictive Fisher information than high-loss and random selectors with the same target within-class mass in every seed (Figure 4). Its temperature-calibrated predictive information is 0.005 ± 0.001 and loss 0.031 ± 0.017, with 95% Student-t intervals across five paired seeds. Accuracy and loss use support-corrected membership

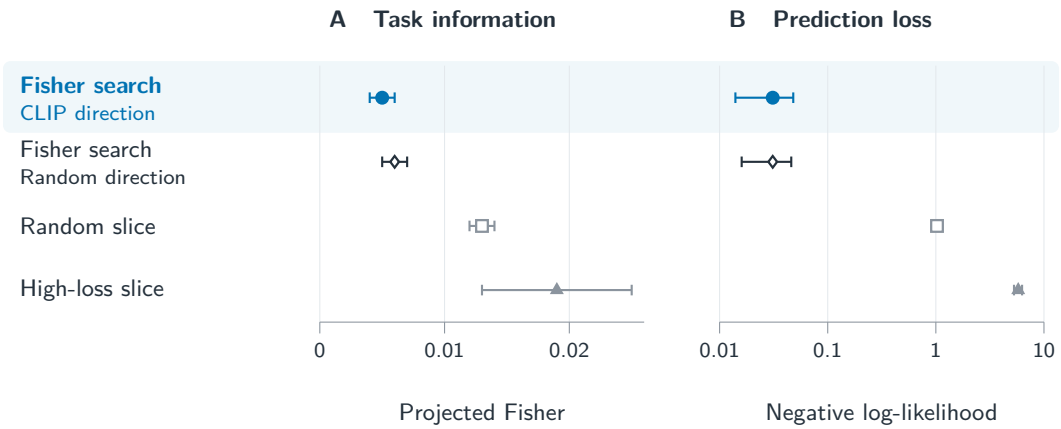


Figure 4: **Discovery finds accurate populations with low predictive information.** Held-out predictive projected Fisher (A) and negative log-likelihood (B; log scale), with means and 95% Student-t intervals across five paired seeds. The highlighted slice has 98.8% accuracy. Discovery uses task labels and a fixed nuisance direction; group annotations remain sealed and the gate is frozen before test evaluation. Predictive moments retain calibration and shrinkage 0.10. The random-direction control attains similar information.

weights without information pooling; Appendix C.1 defines these metrics and overlap conventions. Its top-10% Jaccard overlap with the high-loss slice is 0.001: information search finds a highly accurate population largely distinct from loss-based selection.

The random-direction control attains similar predictive information. The paired CLIP-minus-random interval and both unweighted comparisons with high-loss selection span zero. These unweighted comparisons reevaluate the predictive-selected gates under alternative matrices. The strongest observed separation therefore concerns predictive information, while the incremental contribution of nuisance semantics remains unresolved. Appendix D.4 gives the full table, paired intervals, support diagnostics, and revealed composition.

6 Information-Guided Interventions

The audit supplies objectives for data allocation, training, and acquisition. An observed-source allocation instantiates the Fisher-reserve principle of Mansouri and Arai [2026]: assigning fraction r to an anchor with unregularized projected information at least $\rho_0 I$ establishes the floor $r\rho_0$, under common coordinates and positive-definite nuisance blocks (Appendix E.1). Learning combines empirical task risk or Group DRO with inverse-information tail risk:

$$\min_{w} \ \widehat{R}_{\mathrm{base}}(w) + \lambda_{\mathrm{FG}} \, \mathrm{CVaR}_{\alpha,g}\left[\left\{\lambda_{\min}(\widehat{F}_g^{\mathrm{rid}}(w)) + \varepsilon\right\}^{-1}\right]. \tag{11}$$

Weak projected-direction penalties complement base task risk. Coordinates follow Section 4. Frozen-feature Waterbirds uses validation group labels to select the largest validation certificate among coefficients within two percentage points of the best validation worst-group accuracy.

Table 1: Frozen-feature Waterbirds training. Means with nominal 95% normal-approximation intervals (± 1.96 standard errors) across five seeds; ridge certificates use $\alpha = 0.95$. Validation selection follows Section 6.

Objective	Average accuracy (%)	Worst-group accuracy (%)	Fisher-Glass certificate
ERM	77.2 ± 0.1	55.7 ± 0.3	0.198 ± 0.001
Fisher-Glass	80.4 ± 0.2	56.3 ± 0.3	0.399 ± 0.006
Group DRO	80.9 ± 0.3	67.3 ± 0.3	0.233 ± 0.003
Group DRO + Fisher-Glass	80.6 ± 0.9	69.4 ± 1.3	0.614 ± 0.059

Fisher-Glass approximately doubles ERM's certificate (Table 1); with Group DRO it raises the certificate from 0.233 to 0.614 and mean worst-group accuracy from 67.3% to 69.4%. Stronger penalties can reduce accuracy (Appendix E.3), so selecting an operating point requires task and validation criteria.

Prospective training increases the validation certificate from 0.289 to 0.330, while the controlled background-swap drop is 52.9 points versus 52.6 for ERM. Constrained Fisher-only repair selects the unchanged model for the discovered slice in all five seeds. Post hoc selection increases hidden information but reduces worst-group accuracy by 12.6 points; coefficients coincide with validation-NLL minimization in every seed. Appendix E.3 records the comparisons.

For acquisition, candidate a has modified feature/input x_a^{cf}, assigned label y_a, and current predictive loss $\ell_w(x_a^{\mathrm{cf}}, y_a)$. *Tail-transverse influence* (TTI) multiplies local tail-information gain by this loss, $\mathrm{TTI}(a) = G_\alpha(a)\ell_w(x_a^{\mathrm{cf}}, y_a)$. Gain differentiates the Schur complement along its weakest task direction, weighting environments by inverse-information CVaR sensitivity (Appendix E.4). Waterbirds implements $\widetilde{\mathrm{TTI}}_\alpha(a) = \widetilde{G}_\alpha(a)\ell_w(x_a^{\mathrm{cf}}, y_a)$ with a normalized residual-score surrogate $\widetilde{G}_\alpha$; Appendix E.5 specifies its weighting and allocation conventions.

Waterbirds candidates translate features by a same-label background-group mean difference, scaled by γ_{cf}, using task and background annotations. Training data supply means and the nuisance probe; validation data supply candidates, selected features enter retraining, and evaluation is held out. Budgets count acquisitions relative to original training-set size. At unit feature-shift strength, $\gamma_{\mathrm{cf}} = 1$, and a 1% budget, $\widetilde{\mathrm{TTI}}_\alpha$ yields a mean worst-group accuracy gain of 4.52 points versus 3.27 for high-loss selection; high loss produces the larger mean certificate gain. At this strength, $\widetilde{\mathrm{TTI}}_\alpha$ has the larger mean worst-group accuracy gain at all three evaluated budgets; pure tail gain yields smaller improvements. See Appendix E.5 for construction and budgets.

7 Related Work

Efficient information and classification. Efficient-score theory defines task information after projection away from nuisance tangent spaces [Bickel et al., 1993, Newey, 1990, Tsiatis, 2006]. The Fisher-glass framework of Mansouri and Arai [2026] studies the environmental distribution of that information through geometric tail laws and inverse-information certification. We develop its classification consequences through local testing reductions [Le Cam, 1986, Tsybakov, 2009] under an explicit Bayes-decision separation condition. The operational analysis addresses Schur estimation, ridge effects, and independent evaluation of learned population gates.

Tail risk, robust learning, and gradient geometry. CVaR concentrates objectives on adverse outcomes [Rockafellar and Uryasev, 2000, 2002]; Group DRO applies robust optimization to predictive loss across groups [Sagawa et al., 2020]. Invariant Risk Minimization seeks shared optimal classifiers [Arjovsky et al., 2019], while Fishr regularizes domain-level gradient variances [Rame et al., 2022]. We use inverse projected-information tail risk as a training penalty alongside empirical risk or Group DRO, and evaluate information and predictive performance together.

Hidden stratification and slice discovery. Hidden-stratification research motivates finer population evaluation [Oakden-Rayner et al., 2020]. GEORGE clusters representations to infer subclasses [Sohoni et al., 2020], JTT upweights initially misclassified examples [Liu et al., 2021], and AGRO learns adversarial group assignments using prediction loss [Paranjape et al., 2023]. Our search targets low predictive projected information given task labels and a fixed nuisance direction. Support constraints and an independently evaluated frozen gate make accurately classified, low-loss populations eligible audit targets.

Information-based design and acquisition. E-optimal design emphasizes the smallest information eigenvalue [Pukelsheim, 2006]; BAIT uses Fisher information for neural active learning [Ash et al., 2021]. TTI measures local information gain along a weak nuisance-adjusted task direction, weights environments by tail sensitivity, and multiplies by current task loss. Acquisition experiments assess realized information and accuracy changes after retraining.

8 Limitations

Specification and attribution. The audit depends on task contrast, nuisance space, and representation; nuisance variation must preserve the task distinction. Absolute eigenvalues depend on task scaling; ridge estimates also on nuisance scaling. Standardization supplies a comparison convention, not representation invariance. Predictive Fisher weights confidence; unweighted Schur statistics describe residual feature geometry. Their relation to population efficient information depends on the score model. Centering reverses the prospective ranking in all five seeds, so its original uncentered diagnostic depends partly on slice means. Comparable CLIP and random-direction discovery results and unweighted-control intervals containing zero leave semantic nuisance projection's incremental contribution unresolved. Confidence-matched comparisons and systematic nuisance controls remain necessary. Audits may expose subgroup failures, but treating a local certificate as a deployment guarantee could overlook harms outside the audited populations.

Local theory and audit scope. Efficient information, separation strength, and local radius jointly determine classification bounds, which require regularity and quantitative Bayes-decision separation. Sample size counts the audited environment's observations. Latent-environment pooling also depends on prevalence, shared information, and discovery. Coverage is relative to the declared environment distribution. Finite audits cover included groups or frozen slice rules; small mass and weak nuisance conditioning amplify uncertainty. Implemented Bonferroni-adjusted Gaussian/Wald audits have nominal coverage conditional on fitted model and evaluated family. The reported Wald procedure holds realized normalizations fixed and does not propagate uncertainty from normalizers estimated on evaluation data. Repeated searches must account for certification-data reuse. Empirical sample-size scaling and achievable classification rates remain open.

Development history and empirical breadth. After pilot experiments, the prospective rule was frozen for fresh training seeds on the same Waterbirds benchmark. Repetitions measure variability within this development setting; independent confirmatory datasets are needed. The prospective experiment establishes controlled background sensitivity in the selected group; comparative selection gains require additional matched evaluations. The probe's background attribute defines the intervention; same-type recomposition controls shared construction effects. Evidence centers on MobileNetV3-small image classifiers, TF-IDF text diagnostics, and synthetic constructions. Broader representations and entangled task–nuisance settings need evaluation. Unregularized audits used fresh training runs; checkpoint-specific conclusions concern those replications.

Information, performance, and intervention scope. Information gains do not establish behavioral repair: prospective swap sensitivity remains similar, constrained hidden-slice repair selects unchanged models, and post hoc selection increases hidden information while reducing worst-group accuracy. Prospective models lack accuracy matching in two of five seeds; post hoc choices coincide with validation-NLL minimization. Deployment requires prospective validation against its performance objective. Reserves assume an informative anchor and stated information-accumulation design. TTI measures the current predictor's local sensitivity; finite interventions and retraining are empirical tests. Transfer of annotation-conditioned acquisition feature shifts to realizable label-preserving interventions and other candidate families remains unestablished.

9 Conclusion

We develop classification and auditing methods for evaluating nuisance-projected task information in trained predictors. A local excess-risk lower bound connects efficient information to learning difficulty under explicit decision-separation conditions. Regularization-aware bounds and independent evaluation support audits of predefined groups and learned subpopulations. On Waterbirds, a validation-selected group exhibits a 52.6-point controlled accuracy drop under background replacement, while held-out discovery identifies accurate populations with low estimated projected predictive Fisher information. Training and acquisition experiments characterize the joint response of information and predictive performance. These results support evaluating classifiers through observed errors, specified information targets, and sensitivity to declared nuisance changes.

References

Martin Arjovsky, Léon Bottou, Ishaan Gulrajani, and David Lopez-Paz. Invariant risk minimization. *arXiv preprint arXiv:1907.02893*, 2019. doi: 10.48550/arXiv.1907.02893.

Jordan T. Ash, Surbhi Goel, Akshay Krishnamurthy, and Sham M. Kakade. Gone fishing: Neural active learning with Fisher embeddings. In *Advances in Neural Information Processing Systems*, volume 34, pages 8927–8939. Curran Associates, Inc., 2021. URL https://proceedings.neurips.cc/paper/2021/hash/4afe044911ed2c247005912512ace23b-Abstract.html.

Jean-Yves Audibert and Alexandre B. Tsybakov. Fast learning rates for plug-in classifiers. *The Annals of Statistics*, 35(2):608–633, 2007. doi: 10.1214/009053606000001217.

Peter J. Bickel, Chris A. J. Klaassen, Ya'acov Ritov, and Jon A. Wellner. *Efficient and Adaptive Estimation for Semiparametric Models*. Johns Hopkins University Press, 1993.

Daniel Borkan, Lucas Dixon, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. Nuanced metrics for measuring unintended bias with real data for text classification. In *Companion Proceedings of The 2019 World Wide Web Conference*, pages 491–500. ACM, 2019. doi: 10.1145/3308560.3317593.

D. R. Cox and N. Reid. Parameter orthogonality and approximate conditional inference. *Journal of the Royal Statistical Society: Series B (Methodological)*, 49(1):1–18, 1987. doi: 10.1111/j.2517-6161.1987.tb01422.x.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR, 2017. URL https://proceedings.mlr.press/v70/guo17a.html.

Andrew Howard, Mark Sandler, Grace Chu, Liang-Chieh Chen, Bo Chen, Mingxing Tan, Weijun Wang, Yukun Zhu, Ruoming Pang, Vijay Vasudevan, Quoc V. Le, and Hartwig Adam. Searching for MobileNetV3. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 1314–1324, 2019. doi: 10.1109/ICCV.2019.00140. URL https://openaccess.thecvf.com/content_ICCV_2019/html/Howard_Searching_for_MobileNetV3_ICCV_2019_paper.html.

Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanas Phillips, Irena Gao, Tony Lee, Etienne David, Ian Stavness, Wei Guo, Berton Earnshaw, Imran Haque, Sara M. Beery, Jure Leskovec, Anshul Kundaje, Emma Pierson, Sergey Levine, Chelsea Finn, and Percy Liang. WILDS: A benchmark of in-the-wild distribution shifts. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 5637–5664. PMLR, 2021. URL https://proceedings.mlr.press/v139/koh21a.html.

Lucien Le Cam. *Asymptotic Methods in Statistical Decision Theory*. Springer Series in Statistics. Springer, 1986. doi: 10.1007/978-1-4612-4946-7.

Evan Z. Liu, Behzad Haghgoo, Annie S. Chen, Aditi Raghunathan, Pang Wei Koh, Shiori Sagawa, Percy Liang, and Chelsea Finn. Just train twice: Improving group robustness without training group information. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 6781–6792. PMLR, 2021. URL https://proceedings.mlr.press/v139/liu21f.html.

El Mustapha Mansouri and Keigo Arai. Fisher Glasses: Tail-certified quantum metrology in quenched environments. arXiv preprint arXiv:2607.01085v1, 2026. URL https://arxiv.org/abs/2607.01085v1.

Whitney K. Newey. Semiparametric efficiency bounds. *Journal of Applied Econometrics*, 5(2): 99–135, 1990. doi: 10.1002/jae.3950050202.

Luke Oakden-Rayner, Jared Dunnmon, Gustavo Carneiro, and Christopher Ré. Hidden stratification causes clinically meaningful failures in machine learning for medical imaging. In *Proceedings of the ACM Conference on Health, Inference, and Learning*, pages 151–159. ACM, 2020. doi: 10.1145/3368555.3384468.

Bhargavi Paranjape, Pradeep Dasigi, Vivek Srikumar, Luke Zettlemoyer, and Hannaneh Hajishirzi. AGRO: Adversarial discovery of error-prone groups for robust optimization. In *International*

Conference on Learning Representations, 2023. URL https://openreview.net/forum?id=IrzkT99fDJH.

Friedrich Pukelsheim. *Optimal Design of Experiments*, volume 50 of *Classics in Applied Mathematics*. SIAM, 2006. doi: 10.1137/1.9780898719109.

Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 8748–8763. PMLR, 2021. URL https://proceedings.mlr.press/v139/radford21a.html.

Alexandre Rame, Corentin Dancette, and Matthieu Cord. Fishr: Invariant gradient variances for out-of-distribution generalization. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 18347–18377. PMLR, 2022. URL https://proceedings.mlr.press/v162/rame22a.html.

R. Tyrrell Rockafellar and Stanislav Uryasev. Optimization of conditional value-at-risk. *The Journal of Risk*, 2(3):21–41, 2000. doi: 10.21314/JOR.2000.038.

R. Tyrrell Rockafellar and Stanislav Uryasev. Conditional value-at-risk for general loss distributions. *Journal of Banking & Finance*, 26(7):1443–1471, 2002. doi: 10.1016/S0378-4266(02)00271-6.

Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. ImageNet large scale visual recognition challenge. *International Journal of Computer Vision*, 115: 211–252, 2015. doi: 10.1007/s11263-015-0816-y.

Shiori Sagawa, Pang Wei Koh, Tatsunori B. Hashimoto, and Percy Liang. Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization. In *International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=ryxGuJrFvS.

Nimit S. Sohoni, Jared A. Dunnmon, Geoffrey Angus, Albert Gu, and Christopher Ré. No subclass left behind: Fine-grained robustness in coarse-grained classification problems. In *Advances in Neural Information Processing Systems*, volume 33, pages 19339–19352. Curran Associates, Inc., 2020. URL https://proceedings.neurips.cc/paper/2020/hash/e0688d13958a19e087e123148555e4b4-Abstract.html.

Alexandru Telea. An image inpainting technique based on the fast marching method. *Journal of Graphics Tools*, 9(1):25–36, 2004. doi: 10.1080/10867651.2004.10487596. URL https://doi.org/10.1080/10867651.2004.10487596.

Joel A. Tropp. User-friendly tail bounds for sums of random matrices. *Foundations of Computational Mathematics*, 12(4):389–434, 2012. doi: 10.1007/s10208-011-9099-z.

Anastasios A. Tsiatis. *Semiparametric Theory and Missing Data*. Springer Series in Statistics. Springer, 2006. doi: 10.1007/0-387-37345-4.

Alexandre B. Tsybakov. *Introduction to Nonparametric Estimation*. Springer Series in Statistics. Springer, 2009. doi: 10.1007/b13794.

Catherine Wah, Steve Branson, Peter Welinder, Pietro Perona, and Serge Belongie. The Caltech-UCSD Birds-200-2011 dataset. Technical Report CNS-TR-2011-001, California Institute of Technology, 2011. URL https://www.vision.caltech.edu/datasets/cub_200_2011/.

Bolei Zhou, Agata Lapedriza, Aditya Khosla, Aude Oliva, and Antonio Torralba. Places: A 10 million image database for scene recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(6):1452–1464, 2018. doi: 10.1109/TPAMI.2017.2723009. URL https://doi.org/10.1109/TPAMI.2017.2723009.

A Proofs of the Core Results

Throughout, task coordinates use the common normalization of Section 2, and environments follow the declared distribution Π.

A.1 Projection, Normalization, and Task–Nuisance Tangency

Use the centered, square-integrable scores and blocks $J_{AB,e}$ of Section 2.1. Suppose the finite nuisance scores span $\mathcal{T}_{N,e}$ and $J_{NN,e} \succ 0$. With $\Gamma_e = J_{TN,e} J_{NN,e}^{-1}$,

$$\begin{aligned} \mathbb{E}_{P_e}[(s_{T,e} - \Gamma_e s_{N,e}) s_{N,e}^\top] &= J_{TN,e} - \Gamma_e J_{NN,e} = 0, \\ F_e = \mathbb{E}_{P_e}[s_{T,e}^{\text{eff}} (s_{T,e}^{\text{eff}})^\top] &= J_{TT,e} - J_{TN,e} J_{NN,e}^{-1} J_{NT,e}. \end{aligned} \tag{A.1}$$

Because $\Gamma_e s_{N,e}$ lies componentwise in the nuisance span, orthogonality identifies the residual as the efficient score in Equation (1) [Bickel et al., 1993, Tsiatis, 2006]. Equivalently,

$$u^\top F_e u = \min_{v \in \mathbb{R}^{d_N}} \mathbb{E}_{P_e}[(u^\top s_{T,e} - v^\top s_{N,e})^2], \qquad v^\star = J_{NN,e}^{-1} J_{NT,e} u. \tag{A.2}$$

Hence $F_e \succeq 0$, with $u^\top F_e u = 0$ exactly when $u^\top s_{T,e}$ belongs to the nuisance span. The projection definition extends directly to the closed span of an infinite-dimensional nuisance tangent space.

For separate invertible transformations $\widetilde{s}_T = U s_T$ and $\widetilde{s}_N = V s_N$, the nuisance span is unchanged and

$$\widetilde{s}_T^{\text{eff}} = U s_T^{\text{eff}}, \qquad \widetilde{F} = U F U^\top; \qquad \frac{F}{J_{TT}} = 1 - \rho^2 \quad \text{in the scalar case.} \tag{A.3}$$

Here $\rho = J_{TN}/\sqrt{J_{TT} J_{NN}}$ requires positive raw task and nuisance information. Exact information is invariant to nuisance-basis changes; task congruence preserves rank but generally changes eigenvalues. The scalar surviving fraction is invariant to nonzero task rescaling and invertible nuisance-basis changes. A fixed isotropic nuisance ridge does not share the latter invariance unless its penalty is transformed accordingly (Appendix B).

For the tangency construction in Section 2.2, let $b_e, r_e \in L_0^2(P_e)$ be orthogonal unit vectors, $\mathcal{T}_{N,e} = \operatorname{span}\{b_e\}$, and $a_e = b_e + \delta_e r_e$ with $\delta_e \geq 0$. The projection coefficient is $\langle a_e, b_e \rangle_e / \|b_e\|_e^2 = 1$, so

$$a_e^{\text{eff}} = \delta_e r_e, \qquad F_e = \|\delta_e r_e\|_e^2 = \delta_e^2, \qquad J_{TT,e} = 1 + \delta_e^2, \qquad \frac{F_e}{J_{TT,e}} = \frac{\delta_e^2}{1 + \delta_e^2}. \tag{A.4}$$

This construction realizes the transverse-information geometry of Mansouri and Arai [2026]; coexistence with low prediction risk additionally requires a successful nuisance-based predictor.

A.2 Normal Coordinates and the Inverse-Information Tail

Under Theorem 1, the zero set $\mathcal{Z} = \{\lambda = 0\}$ is a nonempty compact C^2 embedded manifold without boundary, of codimension $c \in \{1, \ldots, d\}$. The nonnegative λ is continuous on the closure of a tube $U_{r_\star}(\mathcal{Z})$ and bounded away from zero outside it. In coordinates $z = s + r\omega$, with $\omega \in S(N_s \mathcal{Z})$, assume uniformly on the unit normal bundle

$$\lambda(s + r\omega) = a(s, \omega) r^{2m} (1 + o(1)), \qquad p(s + r\omega) = b(s, \omega) r^\nu (1 + o(1)), \tag{A.5}$$

where p is the Lebesgue density, $m > 0$, $\nu > -c$, and continuous a, b have positive uniform lower and upper bounds. Put

$$\beta = \frac{c + \nu}{2m}, \qquad C_{\mathcal{Z}} = \frac{1}{c + \nu} \int_{\mathcal{Z}} \int_{S(N_s \mathcal{Z})} b(s, \omega) a(s, \omega)^{-\beta} \, d\sigma_s(\omega) \, d \operatorname{vol}_{\mathcal{Z}}(s). \tag{A.6}$$

The sphere measure σ_s is unnormalized; for $c = 1$, each of the two normal directions has mass one.

Proof of Theorem 1. For any $0 < r_0 < r_\star$, continuity and absence of zeros on the compact annulus $r_0 \leq \operatorname{dist}(z, \mathcal{Z}) \leq r_\star$ give a positive minimum of λ. Together with the bound outside the outer tube, this localizes $\{\lambda < t\}$ to $U_{r_0}(\mathcal{Z})$ for all sufficiently small t. The tubular volume element

is $dz = j(s, r\omega) r^{c-1}\, dr\, d\sigma_s\, d\,\mathrm{vol}_{\mathcal{Z}}$, where $j = 1 + O(r)$ uniformly. Thus, for any $0 < \zeta < 1$, uniformity permits r_0 small enough that

$$\begin{aligned} (1-\zeta) a r^{2m} \le \lambda(s + r\omega) \le (1+\zeta) a r^{2m}, \\ (1-\zeta) b r^{\nu} \le p(s + r\omega) j(s, r\omega) \le (1+\zeta) b r^{\nu} \end{aligned} \qquad (0 < r < r_0). \tag{A.7}$$

For small t, all radial thresholds lie inside this tube. The event contains $r < (t/((1+\zeta)a))^{1/(2m)}$ and is contained in $r < (t/((1-\zeta)a))^{1/(2m)}$. Integrating uses $\int_0^{(t/a)^{1/(2m)}} b r^{c+\nu-1}\, dr = b a^{-\beta} t^{\beta}/(c+\nu)$ and gives

$$(1-\zeta)(1+\zeta)^{-\beta} C_{\mathcal{Z}} \le \frac{\Pr(\lambda < t)}{t^{\beta}} \le (1+\zeta)(1-\zeta)^{-\beta} C_{\mathcal{Z}}. \tag{A.8}$$

Let $t \downarrow 0$ and then $\zeta \downarrow 0$. Compactness and the coefficient bounds ensure $0 < C_{\mathcal{Z}} < \infty$, so $\Pr(\lambda < t) \sim C_{\mathcal{Z}} t^{\beta}$. Since $c \ge 1$, $\mathcal{Z}$ has zero Lebesgue and environment measure; substituting $t = x^{-1}$ proves the inverse tail in the next display. □

The tail-integral identity then yields, for every $p_0 > 0$,

$$\begin{aligned} &\Pr(\lambda^{-1} > x) \sim C_{\mathcal{Z}} x^{-\beta}, \\ &\mathbb{E}[\lambda^{-p_0}] = p_0 \int_0^{\infty} x^{p_0 - 1} \Pr(\lambda^{-1} > x)\, dx < \infty \quad \Longleftrightarrow \quad p_0 < \beta. \end{aligned} \tag{A.9}$$

The bounded-x part is finite; the tail has the convergence criterion of $\int^{\infty} x^{p_0 - \beta - 1}\, dx$. In particular, the inverse mean is finite exactly for $\beta > 1$, with logarithmic divergence at $\beta = 1$.

Scalar branches. For $\lambda = \delta^2$, $\Pr(\delta < t) \sim C_\delta t^{\gamma}$ gives exponent $\gamma/2$. More generally, let $\lambda_K = \sum_{k=1}^{K} w_k \delta_k^2$, $w_k > 0$, where $\delta \in [0, \infty)^K$ has joint density f continuous and positive at the origin from within the orthant. Changing variables gives the scalar portfolio law of Mansouri and Arai [2026]:

$$\begin{aligned} \Pr(\lambda_K < t) &= \frac{t^{K/2}}{\sqrt{\prod_k w_k}} \int_{B_K^+} f\big(\sqrt{t/w_1} u_1, \ldots, \sqrt{t/w_K} u_K\big)\, du \\ &\sim \frac{f(0) v_K}{2^K \sqrt{\prod_k w_k}}\, t^{K/2}. \end{aligned} \tag{A.10}$$

Here B_K^+ is the positive-orthant unit ball and v_K the full ball's volume. Independence gives $f(0) = \prod_k f_k(0+)$; the inverse mean is finite for $K \ge 3$. This calculation concerns additive scalar information or a fixed task direction. A minimum-eigenvalue conclusion instead requires applying Theorem 1 to that eigenvalue.

A.3 Local Classification: Testing Reduction and Tail Aggregation

Let $\mathcal{E}_P(g) = R_P(g) - R_P^{\star}$. Fix a local family $\mathcal{P}_e^{\mathrm{loc}}$ independently of sample size, and define

$$\mathfrak{R}_e(n) = \inf_{\widehat{g}_n} \sup_{P \in \mathcal{P}_e^{\mathrm{loc}}} \mathbb{E}_{S_n \sim P^{\otimes n}}[\mathcal{E}_P(\widehat{g}_n(S_n))], \tag{A.11}$$

where learning rules are measurable and n counts independent observations from environment e; independent learner randomization is included in the expectation. Assume $\lambda_e > 0$ and a unit weakest efficient-information direction u_e. For every $0 < t \le t_{0,e}$, the fixed family contains $P_{e,\pm t}$ with task coordinates $\pm t u_e$ and potentially varying nuisance coordinates. For fixed $K_e, c_e, t_{0,e} > 0$ and $q > 0$, assume both bounds in Equation (5) hold throughout that radius.

Proof of Theorem 2. Fix an admissible t and abbreviate $P_{\pm} = P_{e,\pm t}$. For any learning rule, the decision-separation condition $\mathcal{E}_{P_+}(g) + \mathcal{E}_{P_-}(g) \ge c_e t^q$ implies that choosing the alternative with smaller excess risk makes any testing error cost at least $c_e t^q/2$ under the true alternative. Explicitly,

$$\begin{aligned} &\psi_t(S_n) = + \iff \mathcal{E}_{P_+}(\widehat{g}_n(S_n)) \le \mathcal{E}_{P_-}(\widehat{g}_n(S_n)), \\ &\max_{\sigma \in \{+,-\}} \mathbb{E}_{P_\sigma^{\otimes n}}[\mathcal{E}_{P_\sigma}(\widehat{g}_n)] \ge \frac{c_e t^q}{4} \{P_+^{\otimes n}(\psi_t = -) + P_-^{\otimes n}(\psi_t = +)\}. \end{aligned} \tag{A.12}$$

The standard exponential two-point testing inequality [Tsybakov, 2009, Lemma 2.6], together with product additivity of KL and the assumed KL bound, gives

$$P_+^{\otimes n}(\psi_t = -) + P_-^{\otimes n}(\psi_t = +) \geq \tfrac{1}{2}\exp[-n\,\mathrm{KL}(P_+\|P_-)] \geq \tfrac{1}{2}e^{-nK_e t^2 \lambda_e},$$
$$\mathfrak{R}_e(n) \geq \frac{c_e}{8} t^q e^{-nK_e t^2 \lambda_e}. \tag{A.13}$$

The second line takes the supremum over the fixed family containing both alternatives and then the infimum over learning rules. Choose

$$t_n = \min\{t_{0,e}, (nK_e\lambda_e)^{-1/2}\}, \qquad nK_e\lambda_e t_n^2 \leq 1,$$
$$\mathfrak{R}_e(n) \geq c_0 c_e \min\{t_{0,e}^q, (nK_e\lambda_e)^{-q/2}\}, \qquad c_0 = \frac{e^{-1}}{8}. \tag{A.14}$$

Independent learner randomization has the same law under both alternatives and leaves their KL divergence unchanged, so the argument also covers randomized rules. □

Quadratic decision separation. Suppose the binary alternatives share a feature marginal Q_e, with regression functions $\eta_{\pm,t}(x) = P_{e,\pm t}(Y = 1 \mid X = x)$. Let $D_{e,t} = \{g^\star_{+,t} \neq g^\star_{-,t}\}$ be their Bayes-rule disagreement set. Applying the binary excess-risk identity [Audibert and Tsybakov, 2007, Eq. (1.1)] to the two alternatives gives

$$\inf_g \{\mathcal{E}_{P_{e,+t}}(g) + \mathcal{E}_{P_{e,-t}}(g)\}$$
$$= \int_{D_{e,t}} \min\{|2\eta_{+,t}(x) - 1|, |2\eta_{-,t}(x) - 1|\}\, dQ_e(x). \tag{A.15}$$

Pointwise, a classifier must disagree with at least one Bayes rule on $D_{e,t}$; choosing the smaller margin attains the integral for unrestricted measurable binary classifiers, while it remains a lower bound for a restricted class. If, for every $0 < t \leq t_{0,e}$, a measurable $A_{e,t} \subseteq D_{e,t}$ has $Q_e(A_{e,t}) \geq c_{\mathrm{mass},e}t$ and both margins at least $c_{\mathrm{margin},e}t$, with positive constants independent of t, the integral is at least $c_{\mathrm{mass},e}c_{\mathrm{margin},e}t^2$. Thus $q = 2$ and $c_e = c_{\mathrm{mass},e}c_{\mathrm{margin},e}$ satisfy decision separation; the KL condition is still required separately.

Environmental aggregation. If these assumptions hold for Π-almost every environment and all quantities are measurable, monotonicity and positive homogeneity of CVaR give

$$\mathcal{B}_e(n) = c_e \min\{t_{0,e}^q, (nK_e\lambda_e)^{-q/2}\},$$
$$\mathrm{CVaR}_{\alpha,e\sim\Pi}(\mathfrak{R}_e(n)) \geq c_0\,\mathrm{CVaR}_{\alpha,e\sim\Pi}(\mathcal{B}_e(n)). \tag{A.16}$$

Separation c_e, distinguishability K_e, and radius $t_{0,e}$ remain inside the aggregation. For $q = 2$ and $nK_e\lambda_e \geq t_{0,e}^{-2}$, $\mathcal{B}_e(n) = c_e/(nK_e\lambda_e)$; outside that interior regime the radius truncates the bound.

A.4 CVaR, Information Coverage, and Stabilization

For $0 < \alpha < 1$ and nonnegative Z, the CVaR minimization formula holds for integrable Z [Rockafellar and Uryasev, 2002, Theorem 10]; the threshold-weight argument below gives its dual:

$$\mathrm{CVaR}_\alpha(Z) = \inf_{\tau\in\mathbb{R}} \left\{\tau + \frac{\mathbb{E}[(Z-\tau)_+]}{1-\alpha}\right\} = \sup_{\xi\in\mathcal{W}_\alpha} \mathbb{E}[\xi Z],$$
$$\mathcal{W}_\alpha = \{\xi : 0 \leq \xi \leq (1-\alpha)^{-1},\ \mathbb{E}\xi = 1\}. \tag{A.17}$$

Both values are $+\infty$ if $\mathbb{E}Z = +\infty$. For integrable Z, a maximizing weight equals $(1-\alpha)^{-1}$ above $v_\alpha = \inf\{v : \Pr(Z \leq v) \geq \alpha\}$, zero below it, and fills the remaining mass at the threshold. Such fractional weighting is possible because $\Pr(Z > v_\alpha) \leq 1 - \alpha \leq \Pr(Z \geq v_\alpha)$. The dual gives monotonicity, positive homogeneity, and $v_\alpha \leq \mathrm{CVaR}_\alpha(Z)$. For a finite audit with $\pi_g \geq 0$, $\sum_g \pi_g = 1$,

$$\mathrm{CVaR}_{\alpha,\pi}(Z_g) = \max_{\substack{\sum_g \omega_g = 1 \\ 0\leq\omega_g\leq\pi_g/(1-\alpha)}} \sum_g \omega_g Z_g. \tag{A.18}$$

These normalized tail weights, including threshold atoms, also govern Appendix E.4.

Proof of Proposition 1. Let $\varepsilon > 0$ and $\lambda_e \in [0, \infty)$ be finite almost surely. Write $L_e^{(\varepsilon)} = (\lambda_e + \varepsilon)^{-1}$ and $I_\alpha^{\mathrm{FG}}(\varepsilon)$ for the certificate in Equation (7). If $I_\alpha^{\mathrm{FG}}(\varepsilon) \geq \rho > \varepsilon$, then

$$\begin{aligned} v_\alpha(L^{(\varepsilon)}) \leq \mathrm{CVaR}_\alpha(L^{(\varepsilon)}) = [I_\alpha^{\mathrm{FG}}(\varepsilon)]^{-1} \leq \rho^{-1}, \\ \Pr_{e\sim\Pi}(\lambda_e \geq \rho - \varepsilon) = \Pr_{e\sim\Pi}(L_e^{(\varepsilon)} \leq \rho^{-1}) \geq \alpha. \end{aligned} \tag{A.19}$$

The last step is the quantile definition and includes equality at the information threshold. □

Since $0 < L^{(\varepsilon)} \leq \varepsilon^{-1}$, always $I_\alpha^{\mathrm{FG}}(\varepsilon) \geq \varepsilon$; only $\rho > \varepsilon$ certifies a strictly positive floor. With four equally weighted groups and $\alpha = 0.95$, one group already contains the entire upper-tail mass, so

$$\mathrm{CVaR}_{0.95}(L^{(\varepsilon)}) = \max_g L_g^{(\varepsilon)}, \qquad I_{0.95}^{\mathrm{FG}}(\varepsilon) = \min_g \lambda_g + \varepsilon. \tag{A.20}$$

Exact limit and positivity. Set $L^{(0)} = \lambda^{-1}$, with $1/0 = +\infty$. As $\varepsilon \downarrow 0$, $L^{(\varepsilon)} \uparrow L^{(0)}$, and monotone convergence applies to $\mathbb{E}[\xi L^{(\varepsilon)}]$ for every feasible ξ. The supremum over ξ commutes with the increasing supremum along any sequence $\varepsilon_k \downarrow 0$, giving

$$\begin{aligned} \mathrm{CVaR}_\alpha(L^{(\varepsilon)}) \uparrow \mathrm{CVaR}_\alpha(\lambda^{-1}), \\ I_\alpha^{\mathrm{FG}}(\varepsilon) \downarrow I_\alpha^{\mathrm{FG}}(0) := \left[\mathrm{CVaR}_\alpha(\lambda^{-1})\right]^{-1}, \qquad 1/(+\infty) = 0. \end{aligned} \tag{A.21}$$

For nonnegative Z, the feasible weight $\xi = 1$ and variational choice $\tau = 0$ imply

$$\mathbb{E}Z \leq \mathrm{CVaR}_\alpha(Z) \leq \frac{\mathbb{E}Z}{1-\alpha}, \qquad I_\alpha^{\mathrm{FG}}(0) > 0 \iff \mathbb{E}[\lambda^{-1}] < \infty. \tag{A.22}$$

Under the geometric tail assumptions, this is equivalent to $\beta = (c + \nu)/(2m) > 1$; the exact certificate vanishes for $\beta \leq 1$. Information coverage is a population statement under Π, separate from the confidence level of an estimated lower certificate.

B Estimation and Certification

The bounds condition on a predictor, coordinates, calibration, and slices fixed independently of evaluation data. Evaluation-fitted normalization requires additional error control. Expanded implementation details and numerical records are preserved in the artifact (see the artifact availability statement), under `docs/extended_appendix_details/`.

B.1 Operational Models and Matrix Targets

For evaluation distribution P_g and calibrated logits ℓ_w, declare task and nuisance perturbations

$$\ell_{\theta,\mu}(x) = \ell_w(x) + \Phi_T(x)\theta + \Phi_N(x)\mu, \qquad D = (\Phi_T\ \Phi_N), \quad d = d_T + d_N. \tag{B.1}$$

With $p = \mathrm{softmax}(\ell_w)$, $H = \mathrm{diag}(p) - pp^\top$, and normalized feature coordinates $v = (v_T^\top\ v_N^\top)^\top$, distinguish

$$\begin{aligned} M^{\mathrm{pred}} = D^\top H D, \qquad M^{\mathrm{obs}} = ss^\top, \quad s = D^\top(\mathbf{e}_Y - p), \\ M^{\mathrm{feat}} = vv^\top, \qquad J_g = \mathbb{E}_{P_g}[M]. \end{aligned} \tag{B.2}$$

Predictive Fisher averages likelihood-score outer products under the fitted conditional label distribution; its scores have conditional mean zero even when features do not. Its Schur complement is efficient information for this operational model. Observed-label moments equal predictive Fisher under conditional correct specification, $P_g(Y = \cdot \mid X) = p(X)$; otherwise they are second moments, whose centered covariance subtracts $\mathbb{E}[s]\mathbb{E}[s]^\top$. Feature geometry instead satisfies $J_g = \Sigma_g + \mu_g\mu_g^\top$, with $\mu_g = \mathbb{E}[v]$ and $\Sigma_g = \mathrm{Cov}(v)$. The prospective centered comparison uses within-slice covariance Σ_g with divisor n_g and unchanged global scales, removing the mean contribution without rescaling.

Prospective selection uses uncentered feature–Schur fractions; Appendix D.2 compares centered and predictive targets. Normalizing $p(1-p)$ weights within a group rescales the effective nuisance ridge and is distinct from the predictive moment. Hidden discovery uses the unnormalized predictive factor and retains same-class shrinkage $\eta = 0.10$. Removing ridge targets the unregularized Schur complement of the declared matrix, including any retained shrinkage. Normalization conventions and nuisance-specification checks appear in Appendices D and D.5.

B.2 Ridge Regularization and the Unregularized Target

For $J = \left(\begin{smallmatrix} A & B \\ B^\top & C \end{smallmatrix}\right) \succeq 0$, $C \succ 0$, and $\lambda_{\mathrm{rid}} > 0$, define

$$S_0(J) = A - BC^{-1}B^\top, \qquad S_{\mathrm{rid}}(J) = A - B(C + \lambda_{\mathrm{rid}} I)^{-1} B^\top. \tag{B.3}$$

Both are positive semidefinite. The resolvent identity, whose inverse factors commute, gives

$$\begin{aligned} \Delta_{\mathrm{rid}}(J) &:= S_{\mathrm{rid}}(J) - S_0(J) = \lambda_{\mathrm{rid}} BC^{-1}(C + \lambda_{\mathrm{rid}} I)^{-1} B^\top \succeq 0, \\ \|\Delta_{\mathrm{rid}}(J)\|_{\mathrm{op}} &\le \frac{\lambda_{\mathrm{rid}} M^2}{\kappa(\kappa + \lambda_{\mathrm{rid}})} \quad \text{if } C \succeq \kappa I,\ \|B\|_{\mathrm{op}} \le M. \end{aligned} \tag{B.4}$$

Conditioning is essential: scalar blocks $A = 1$, $B = \sqrt{c}$, $C = c > 0$ give a positive-semidefinite joint matrix with $S_0 = 0$ and $S_{\mathrm{rid}} = \lambda_{\mathrm{rid}}/(c + \lambda_{\mathrm{rid}}) \to 1$ as $c/\lambda_{\mathrm{rid}} \downarrow 0$. Comparisons therefore retain the declared standardized nuisance coordinates and ridge value.

B.3 Error Allowances and Lower Certificates

Proposition B.1 (Simultaneous lower certificate)**.** *Fix G groups and weights $\pi_g > 0$ summing to one. Within each group, let M_{gi} be i.i.d. with mean J_g, $0 \preceq M_{gi} \preceq R_g I_d$, for specified finite R_g, and $\widehat{J}_g = n_g^{-1} \sum_i M_{gi}$. For $0 < \delta < 1$, set*

$$\begin{aligned} u_\delta &= \log(2dG/\delta), \qquad t_g = R_g \left(\sqrt{2u_\delta/n_g} + 2u_\delta/(3n_g) \right), \\ \mathcal{E}_\delta &= \bigcap_g \{ \|\widehat{J}_g - J_g\|_{\mathrm{op}} \le t_g \}, \qquad \Pr(\mathcal{E}_\delta) \ge 1 - \delta. \end{aligned} \tag{B.5}$$

Assume each $C_g \succ 0$. Define observable bounds and allowances by

$$\begin{aligned} \overline{M}_g &= \|\widehat{B}_g\|_{\mathrm{op}} + t_g, & \underline{\kappa}_g &= \lambda_{\min}(\widehat{C}_g) - t_g, \\ h_{\mathrm{rid}}(t, M) &= t + (2Mt + t^2)/\lambda_{\mathrm{rid}} + M^2 t/\lambda_{\mathrm{rid}}^2, & a_g &= h_{\mathrm{rid}}(t_g, \overline{M}_g), \\ d_g &= \frac{\lambda_{\mathrm{rid}} \overline{M}_g^2}{\underline{\kappa}_g(\underline{\kappa}_g + \lambda_{\mathrm{rid}})} & & \text{when all } \underline{\kappa}_g > 0. \end{aligned} \tag{B.6}$$

Issue (10) *with audit distribution π only when these conditioning margins are positive; otherwise record conditioning failure. Its probability of exceeding the unregularized population certificate is at most δ. On $\mathcal{E}_\delta$, an issued level $\underline{I}_{\alpha,\pi} \ge \rho > \varepsilon$ also implies $\sum_{g:\lambda_{\min}(S_0(J_g)) \ge \rho - \varepsilon} \pi_g \ge \alpha$.*

Proof. For $X_{gi} = M_{gi} - J_g$, $\|X_{gi}\|_{\mathrm{op}} \le R_g$ and $\mathbb{E}[X_{gi}^2] = \mathbb{E}[M_{gi}^2] - J_g^2 \preceq R_g^2 I$. Matrix Bernstein [Tropp, 2012, Theorem 1.4] with the displayed t_g gives failure at most δ/G per group; the union bound allows dependence between group estimates. On this event, block extraction and Weyl give $\|B_g\|_{\mathrm{op}} \le \overline{M}_g$ and $C_g \succeq \underline{\kappa}_g I$, so (B.4) bounds the ridge gap by d_g.

For the estimation error, write $Q = (C + \lambda_{\mathrm{rid}} I)^{-1}$, $\widehat{Q} = (\widehat{C} + \lambda_{\mathrm{rid}} I)^{-1}$, and $\Delta B = \widehat{B} - B$. Positive semidefiniteness and the resolvent identity give $\|Q\|, \|\widehat{Q}\| \le \lambda_{\mathrm{rid}}^{-1}$ and $\|\widehat{Q} - Q\| \le t/\lambda_{\mathrm{rid}}^2$. Decompose the quadratic-block difference as $\Delta B \widehat{Q} \widehat{B}^\top + B\widehat{Q}\Delta B^\top + B(\widehat{Q} - Q)B^\top$. If $\|B\| \le M$, its three norms are bounded by $t(M + t)/\lambda_{\mathrm{rid}}$, $Mt/\lambda_{\mathrm{rid}}$, and $M^2 t/\lambda_{\mathrm{rid}}^2$. Adding the task-block error t proves the bound $h_{\mathrm{rid}}(t, M)$ and hence a_g.

Writing $\lambda_g^0 = \lambda_{\min}(S_0(J_g))$, Weyl now yields

$$\lambda_g^0 \ge \lambda_{\min}(S_{\mathrm{rid}}(J_g)) - d_g \ge \lambda_{\min}(S_{\mathrm{rid}}(\widehat{J}_g)) - a_g - d_g. \tag{B.7}$$

Since $S_0(J_g) \succeq 0$, clipping in (10) preserves $\underline{\lambda}_g \le \lambda_g^0$. Thus monotonicity of CVaR and reciprocation give

$$\underline{I}_{\alpha,\pi} \le I_{\alpha,\pi}^0 := \left[\mathrm{CVaR}_{\alpha, g \sim \pi}((\lambda_g^0 + \varepsilon)^{-1}) \right]^{-1}. \tag{B.8}$$

The information-coverage conclusion follows from Appendix A.4. □

Confidence concerns held-out observations; coverage concerns the fixed weights π. A finite predeclared nuisance family can be included in the union bound over group–specification pairs; the infimum of certificates preserves coverage separately for each specification. Appendix C.4 supplies matrix allowances for frozen soft gates, including random normalization, pooling, and support correction, to which the same deterministic argument applies.

B.4 Implemented Audit and Numerical Comparison

The reported scalar audit uses 10^5 Gaussian/Wald moment draws per environment at *nominal* 95% confidence, separately from the finite-sample proposition. For $u_i = (z_i^2, z_i b_i, b_i^2)^\top$ and normalized weights w_i, it draws $m^* \sim N(\widehat{m}, \widehat{\Sigma}_{\widehat{m}})$, where

$$\widehat{m} = \sum_i w_i u_i, \qquad \widehat{\Sigma}_{\widehat{m}} = \frac{\sum_i w_i^2 (u_i - \widehat{m})(u_i - \widehat{m})^\top}{1 - \sum_i w_i^2}. \tag{B.9}$$

Uniform weights give sample covariance divided by n. Simulation holds realized coordinates and weights fixed and does not propagate evaluation-fitted normalization error. Bonferroni tail $0.05/(3G)$ covers nuisance conditioning, ridge lower information, and an upper ridge gap within a run ($G = 4$ Waterbirds; 18 CivilComments).

Let $C_g^L, L_g^{\mathrm{rid}}, L_g^0$ be lower quantiles clipped to their nonnegative empirical ranges, and U_g^Δ the upper gap quantile enlarged to at least its plug-in value. The implemented lower information is

$$\underline{\lambda}_g^{\mathrm{W}} = \min\{[L_g^{\mathrm{rid}} - U_g^\Delta]_+, L_g^0, \widehat{F}_g^0\}. \tag{B.10}$$

The exact caps only decrease the decomposed value. Invalid draw denominators ($\le 10^{-12}$) give infinite gaps; otherwise the closed-form gap (B.4) is used even for non-PSD Gaussian draws. An infinite upper gap or $C_g^L \le 10^{-12}$ forces zero and an uncertifiable status; support, empirical-PSD, and identity checks can also prevent certification. Settings are $\lambda_{\mathrm{rid}} = 10^{-3}$, $\varepsilon = 10^{-8}$, $\alpha = 0.95$, uniform environment weights, and fractional threshold-atom mass in CVaR. Quantile rules and numerical tolerances are retained in the executable specification.

Table B.1: Same-cohort summaries of 75 default-ridge held-out audits. Conditioning/gap parentheses are within-run nominal Wald lower/upper values; entries are extrema across runs, not joint confidence statements across seeds or methods. Inflation is ridge minus exact plug-in certificate.

		Conditioning and ridge gap		Certificates		
Scope	Audits	Minimum J_{NN}	Maximum group gap	Maximum inflation	Minimum exact	Minimum Wald lower
CivilComments diagnostics	5	1.456 (1.417)	0.00116 (0.00133)	0.00018	0.4970	0.4618
Waterbirds diagnostics	5	0.211 (0.178)	0.00125 (0.00166)	0.00071	0.1639	0.1504
Frozen-feature training	30	0.227 (0.192)	0.00144 (0.00199)	0.00062	0.1972	0.1814
End-to-end training	35	0.349 (0.349)	0.00252 (0.00300)	0.00140	0.2185	0.2004
All audits	75	0.211 (0.178)	0.00252 (0.00300)	0.00140	0.1639	0.1504

Every run passed conditioning and satisfied $\underline{I}_{0.95}^{\mathrm{W}} \le \widehat{I}_{0.95}^0 \le \widehat{I}_{0.95}^{\mathrm{rid}}$; Waterbirds diagnostic values are 0.1504, 0.1639, and 0.1646. The first ordering is enforced by the caps and does not establish confidence coverage. Training audits use fresh models: original training runs lacked the retained environment blocks needed for checkpoint-specific auditing.

The frozen-feature FG and end-to-end rows fix $\lambda_{\mathrm{FG}} = 0.03$ and 10^{-4}, respectively. Frozen Group DRO + FG selects the largest validation certificate within two percentage points of the best validation worst-group accuracy among $\{0.1, 0.3\}$; seeds 17–21 select $(0.3, 0.1, 0.3, 0.3, 0.3)$. Repeated frozen-feature training/audit runs at ridges $10^{-4}, 10^{-3}, 10^{-2}$ give maximum group gaps 0.0001443, 0.0014374, 0.0138286 and Wald upper values 0.0001996, 0.0019880, 0.0191165.

Table B.2: Separate fresh training audits at manuscript-aligned settings: descriptive five-seed means. These differ from full-grid selected runs in the 75-audit extrema.

Training setting	Ridge plug-in	Exact plug-in	Wald lower
Frozen-feature Fisher-Glass	0.4002	0.3999	0.3663
Frozen-feature Group DRO + FG	0.6254	0.6252	0.5664
End-to-end Group DRO + FG	0.5447	0.5442	0.4890

C Hidden-Slice Discovery

This is the authoritative hidden-discovery protocol for Section 4.3; Appendix D.4 reports outcomes. A frozen per-example gate transfers to evaluation data, where normalization and declared support correction are recomputed. Literal prompts, model revision, and execution mechanics remain in the artifact's expanded record (see the artifact availability statement).

C.1 Gate Family and Support Constraints

For frozen standardized features $\psi(x)$ and task class c, define

$$h_{\phi,c}(x) = \sigma(a_c^\top \psi(x) + b_c), \qquad q_i^0 = \frac{h_{\phi,c}(x_i)}{\sum_{j\in\mathcal{D}_c} h_{\phi,c}(x_j)}, \quad \sigma(t) = (1+e^{-t})^{-1},$$

$$m_{\text{nom}} = \tau_{\text{slice}} n_c, \qquad m_{\text{ESS}} = \min\{n_c, \max\{1, 0.8 m_{\text{nom}}\}\}, \qquad b_{\max} = \min\{1, 2/m_{\text{nom}}\}.$$

Here $n_c = |\mathcal{D}_c|$ and $\tau_{\text{slice}} = 0.10$ is within-class mass. Discovery calibrates b_c so that mean membership equals τ_{slice}; the selected checkpoint's intercept is frozen at evaluation, without score re-centering or mass recalibration. The cap denominator is nominal slice size, not the ESS floor.

Support correction projects q^0 onto the capped simplex only if its maximum exceeds $b_{\max}$. Writing the result as q^{cap} and $u_i = 1/n_c$, it takes $\widetilde{q} = u + t(q^{\text{cap}} - u)$ with the largest $t \in [0,1]$ satisfying $\sum_i \widetilde{q}_i^2 \le 1/m_{\text{ESS}}$. Thus $\widetilde{q} \ge 0$, $\sum_i \widetilde{q}_i = 1$, $\max_i \widetilde{q}_i \le b_{\max}$, and $n_{\text{eff}} = (\sum_i \widetilde{q}_i^2)^{-1} \ge m_{\text{ESS}}$. These constraints do not imply nuisance conditioning. Both discovery and evaluation use their own class sample size, so $\widetilde{q}$ depends on the evaluation batch. Reported mass is mean *raw* membership within the selected class; ESS uses $\widetilde{q}$.

Metrics and overlap. Accuracy, uncalibrated binary NLL (temperature one), and post-hoc group composition are averages under $\widetilde{q}$, without information pooling; predicted labels threshold the head logit at zero. Calibrated NLL is stored separately. Fisher and high-loss gates share these conventions. Random controls weight $k = \lceil \tau_{\text{slice}} n_c \rceil$ sampled rows of the Fisher-selected class uniformly; the constant gate uses $1/n_c$. Unweighted and centered-unweighted controls reevaluate the predictive-Fisher and high-loss gates without rediscovery.

Jaccard overlap uses the $k = \max\{1, \lceil \tau_{\text{slice}} n_c \rceil\}$ largest raw memberships in each selected class, intersected on original row identities; different classes give zero. All ten archived Fisher/high-loss cutoffs are untied. Fold stability applies the same top-count rule on common same-class training-calibration reference rows, averaging fold-pair overlaps within each class, then weighting class averages by selected-fold counts; classes with fewer than two selected gates are excluded. Neither comparison uses positive-membership supports or corrected-weight ranks.

C.2 Discovery Objective and Settings

With the declared matrix M_i and same-class pool $\widehat{J}_{\text{pool},c} = n_c^{-1} \sum_{i\in\mathcal{D}_c} M_i$, search each class and select its best feasible checkpoint/restart by discovery utility:

$$\widehat{J}_{\phi,c}^{(\eta)} = (1-\eta) \sum_i \widetilde{q}_i(\phi) M_i + \eta \widehat{J}_{\text{pool},c},$$
$$(c^\star, \phi^\star) \in \arg\max_{c,\phi} \left\{ \lambda_{\min}\left(S_{\lambda_{\text{rid}}}(\widehat{J}_{\phi,c}^{(\eta)}) \right) + \varepsilon \right\}^{-1}, \qquad 0 \le \eta < 1. \tag{C.1}$$

The operational task/nuisance blocks are scalar. The optimizer uses $-\log(\widehat{F}^{\text{rid}} + \varepsilon)$ minus ten times squared relative mass error; the log preserves ranking at fixed mass. The intercept is recalibrated

while AdamW updates gate weights, and a constant gate is an eligible baseline. High-loss search uses the same family, mass, and support rules with weighted prediction-loss utility.

Table C.1: Fixed hidden-discovery settings; sensitivity grids are declared before metadata reveal.

Component	Specification
Representation/nuisance	576 frozen ImageNet MobileNetV3-small features; fixed CLIP ViT-B/32 water-minus-land text direction.
Head and calibration	1,200 steps, learning rate 0.01, weight decay 10^{-4}; temperature minimizes calibration NLL over 400 log-spaced values in $[0.05, 20]$.
Gate optimization	Ten restarts; 1,000 AdamW steps, learning rate 0.03, weight decay 10^{-4}; Gaussian initial weights with standard deviation 0.02.
Matrix and stabilization	$M_i = p_i(1-p_i)v_i v_i^\top$, calibrated p_i, normalized task/nuisance v_i; $\eta = 0.10$, $\lambda_{\mathrm{rid}} = 10^{-3}$, $\varepsilon = 10^{-8}$.
Replication/sensitivity	Seeds 17–21; five validation folds; mass and shrinkage grids each $\{0.05, 0.10, 0.20\}$, headline 0.10.

The predictive factor $p_i(1 - p_i)$ is not normalized away. Shrinkage mixes the complete matrix before the Schur complement and is retained in every evaluation; its uniform pool uses the same class in the currently evaluated sample. Centered-unweighted controls additionally estimate weighted coordinate means. For CLIP, normalize each of five prompt embeddings per concept, average and normalize each concept, then normalize the water-minus-land difference; its dot product with the normalized image embedding is the raw nuisance score. No background or group annotation enters this direction.

C.3 Fitting and Independent Evaluation

Task labels and official splits are available throughout; background and four-way group metadata stay sealed until all seeds finish certification. Table C.2 fixes the fitting/evaluation order. All five primary

Table C.2: Fitting and evaluation sequence. Evaluation weights and pools are recomputed under the frozen rule.

Stage	Fitted or computed	Carried forward
Training/calibration	Label-stratified 85/15 split: fit head and feature/task/nuisance normalizers on 85%; fit temperature on 15%. Raw CLIP scores first use official-training normalization, followed by fit-subset normalization.	Fixed predictor, coordinates, nuisance direction, temperature.
Validation folds	Five shared task-stratified assignments across selectors; discover on complementary folds, evaluate transfer on held-out folds; assess stability on training-calibration rows.	Protocol assessment; fold gates are not ensembled.
Full validation	Refit under the fixed protocol; select class, gate, and restart by discovery utility.	Freeze gate weights, calibrated intercept, selected class, support settings.
Test	Apply frozen memberships; recompute admissible weights, same-class pool, and metrics before metadata reveal.	Held-out results, then post-hoc interpretation.

Fisher gates have inactive archived test support correction. For high-loss seed 19 it raises ESS from 94.48 to 102.72. Frozen membership therefore does not imply frozen numerical weights or inactive correction on bootstrap samples.

C.4 Held-Out Targets and Uncertainty

Condition on all fitting data/objects and a selected-class distribution P_c. For its frozen gate h and matrix map M, the unshrunk target and empirical ratio are

$$J_h = \frac{A}{p}, \quad A = \mathbb{E}_{P_c}[hM], \quad p = \mathbb{E}_{P_c}[h] > 0; \qquad \widehat{J}_h = \frac{\widehat{A}}{\widehat{p}}, \quad (\widehat{A}, \widehat{p}) = \frac{1}{n}\sum_{i=1}^{n}(h_i M_i, h_i). \tag{C.2}$$

Proposition C.1 (Frozen-gate ratio bound)**.** *For i.i.d. evaluation observations from P_c, assume $0 \le h \le 1$ and $0 \preceq M \preceq M_{\max} I_d$, with finite $M_{\max}$ specified independently of evaluation data. Given $0 < \delta < 1$, let*

$$u = \log(4d/\delta), \qquad r_A = M_{\max}\left(\sqrt{2u/n} + 2u/(3n)\right), \qquad r_p = \sqrt{\log(4/\delta)/(2n)}.$$

With conditional probability at least $1-\delta$, simultaneously $\|\widehat{A} - A\|_{\mathrm{op}} \le r_A$ and $|\widehat{p} - p| \le r_p$. Whenever $\underline{p} = \widehat{p} - r_p > 0$ on this event,

$$\|\widehat{J}_h - J_h\|_{\mathrm{op}} \le r_J := \frac{r_A + \|\widehat{J}_h\|_{\mathrm{op}} r_p}{\widehat{p} - r_p}. \tag{C.3}$$

Proof. The centered matrices $h_i M_i - A$ have norm at most $M_{\max}$ and variance bounded by $M_{\max}^2 I_d$. Matrix Bernstein [Tropp, 2012, Theorem 1.4] gives the numerator event with failure at most $\delta/2$; Hoeffding gives the denominator event with failure at most $\delta/2$. Their union bound and

$$\widehat{J}_h - J_h = \frac{(\widehat{A} - A) - \widehat{J}_h(\widehat{p} - p)}{p}, \qquad p \ge \widehat{p} - r_p > 0,$$

prove the observable bound. The shared random denominator is therefore explicitly controlled. □

Pooling and support correction. Set $J_{\mathrm{pool},c} = \mathbb{E}_{P_c}[M]$ and $J_h^{(\eta)} = (1-\eta)J_h + \eta J_{\mathrm{pool},c}$. Allocate $\delta_{\mathrm{ratio}} + \delta_{\mathrm{pool}} = \delta$, using the former in Proposition C.1. For the evaluation mean $\widehat{J}_{\mathrm{pool},c} = n^{-1}\sum_i M_i$, matrix Bernstein supplies

$$r_{\mathrm{pool}} = M_{\max}\left[\sqrt{\frac{2\log(2d/\delta_{\mathrm{pool}})}{n}} + \frac{2\log(2d/\delta_{\mathrm{pool}})}{3n}\right].$$

Writing $\widetilde{J}_h = \sum_i \widetilde{q}_i M_i$, the triangle inequality gives $\|\widetilde{J}_h - \widehat{J}_h\|_{\mathrm{op}} \le M_{\max}\sum_i |\widetilde{q}_i - q_i^0| =: r_{\mathrm{proj}}$. Hence, on the simultaneous event and when $\underline{p} > 0$,

$$\|(1-\eta)\widetilde{J}_h + \eta\widehat{J}_{\mathrm{pool},c} - J_h^{(\eta)}\|_{\mathrm{op}} \le (1-\eta)(r_J + r_{\mathrm{proj}}) + \eta r_{\mathrm{pool}}. \tag{C.4}$$

The union bound permits dependence between pool and ratio estimates; r_{proj} covers batch-dependent support correction. If the pool is fixed from fitting data, it belongs to the conditional target and needs no evaluation-error term. Appendix B.3 converts the matrix allowance into a Schur/eigenvalue lower certificate, requiring positive lower membership mass and nuisance-conditioning margins. The target retains shrinkage; evaluation-estimated centering additionally requires first-moment control.

Reported uncertainty. Each of 2,000 within-selected-class bootstrap replicates resamples aligned membership, task, nuisance, logit, and label rows, recomputing normalized/support-corrected weights and the same-class pool. Gate coefficients/intercept, coordinate normalizers, and temperature stay frozen. The 95% interval is the 2.5th/97.5th percentile of predictive projected information (linear quantile interpolation). Across five paired seeds, form $\Delta_s = T_s^{\mathrm{gate}} - T_s^{\mathrm{control}}$ first, then use their sample standard deviation s_Δ to report $\overline{\Delta} \pm t_{0.975,4} s_\Delta/\sqrt{5}$; single-procedure intervals use the same rule on its five values. Bootstrap uncertainty is conditional on fitted objects; seed intervals reflect fitting variation on the shared benchmark. Neither is a computation or validation of the finite-sample radii above.

D Experimental Protocols and Additional Results

D.1 Experimental Quantities and Aggregation

Table D.1 identifies the operational targets; Appendix B.1 defines their normalization. Real-data audits use nuisance ridge 10^{-3} and stabilization 10^{-8}. Synthetic stabilization is $\varepsilon = 10^{-12}$. Certificates use $\alpha = 0.95$. Primary predefined-group coordinates use training moments; prospective normalization is specified below.

Table D.1: Diagnostic targets and replication units.

Experiment	Representation and reported quantity	Seeds
Prospective Waterbirds	Fine-tuned MobileNetV3-small; background-probe surviving fraction	22–26
Synthetic	Normalized task/nuisance coordinates; inverse-information tail and certificate	11–15
Hidden Waterbirds slices	Frozen MobileNetV3-small and fixed CLIP direction; shrunken predictive information	17–21
CivilComments groups	TF-IDF logistic classifier; feature–Schur certificate	13–17
Waterbirds groups	Frozen MobileNetV3-small and logistic fitting; feature–Schur certificate	17–21

Uncertainty. Prospective, synthetic, predefined-group, and frozen-feature training summaries report nominal 95% normal-approximation half-widths $1.96s/\sqrt{5}$, with sample standard deviation s. Hidden-slice and group-free comparisons use five-seed Student-t intervals; contrasts are paired within seed. The end-to-end frontier reports standard deviations, and the conditional Gaussian/Wald audit is separate (Appendix B.4). Hidden-gate bootstrap intervals condition on the frozen gate and recompute evaluation weights (Appendix C). Acquisition retains means without recovered per-seed intervals.

Seeds describe fitting/partition or control variation on fixed benchmark splits, not independent datasets. The frozen Waterbirds classifier and primary diagnostic are deterministic after feature extraction and logistic fitting; only its random controls and reserve anchor vary. Prospective intervals use training seeds: 23,176 composites share 5,794 birds, with two donors per bird and arm. Accuracy changes are percentage points; information retains its original units.

D.2 Prospective Nuisance Intervention

Training and frozen selection. ImageNet-initialized MobileNetV3-small [Russakovsky et al., 2015, Howard et al., 2019] is fully fine-tuned on Waterbirds [Sagawa et al., 2020] for seeds 22–26: 50 epochs, batch size 32, learning rate 10^{-4}, and bf16 on an H100. A validation-fitted temperature and class-balanced logistic background probe use the pooled representation. Five probe folds are stratified by validation group (3,000 maximum iterations); features use validation standardization. Validation nuisance scores are out of fold and use validation moments; test scores use the full-validation probe and test moments. Task logits use training moments. Test normalization never selects the slice. The rule is

$$\mathcal{G}_{\text{eligible}} = \{g : \mathrm{Acc}_v(g) \geq \mathrm{Acc}_v, \quad \mathrm{NLL}_v(g) \leq \mathrm{median}_h\, \mathrm{NLL}_v(h)\},$$
$$g^\star \in \operatorname*{arg\,min}_{g \in \mathcal{G}_{\text{eligible}}} \lambda_{\min}(\widehat{F}_g^{\text{rid}})/\widehat{J}_{TT,g}. \tag{D.1}$$

Here v denotes validation and eligibility uses uncalibrated NLL. Temperature scaling [Guo et al., 2017] minimizes full-validation binary NLL over 400 log-spaced temperatures in $[0.05, 20]$, choosing the smaller temperature at ties. The checkpoint, temperature, probe, slice and audit constants are frozen before scoring; checkpoint mismatches abort evaluation. Table D.2 shows that LL (landbird/land) and WW (waterbird/water) qualify; WW is selected in all five seeds.

Foreground-preserving construction. Waterbirds combines CUB birds [Wah et al., 2011] with Places backgrounds [Zhou et al., 2018] under soft CUB mattes. For each bird, we retain background pixels verbatim outside the dilated foreground and fill the covered region using Telea inpainting [Telea, 2004] (radius 5, dilation 2). Its extracted bird layer is recomposed onto training-image

Table D.2: Validation selection. Means ± nominal 95% normal-approximation half-widths across seeds 22–26; accuracy is a fraction. Information uses uncentered feature moments.

(a) Performance and eligibility

Group	n	Accuracy	NLL	Eligible
Landbird/land	467	0.990 ± 0.005	0.038 ± 0.016	Yes
Landbird/water	466	0.703 ± 0.019	1.670 ± 0.156	No
Waterbird/land	133	0.390 ± 0.021	7.718 ± 0.800	No
Waterbird/water	133	0.920 ± 0.013	0.638 ± 0.038	Yes

(b) Information

Group	Raw information	Projected information	Surviving fraction
Landbird/land	0.737 ± 0.046	0.436 ± 0.045	0.590 ± 0.027
Landbird/water	0.794 ± 0.032	0.289 ± 0.014	0.365 ± 0.026
Waterbird/land	0.973 ± 0.033	0.764 ± 0.031	0.786 ± 0.036
Waterbird/water	2.063 ± 0.054	0.511 ± 0.021	0.248 ± 0.015

donors of the same or opposite background type. Donors have foreground fraction at most 0.08; a deterministic sequence tries four, taking the first with visible inpainted area at most 0.03, or the least-inpainted of four. The same-type arm undergoes identical recomposition.

Both arms use the same foreground layer, matte, label and source birds. Opaque foreground pixels are unchanged before JPEG encoding (quality 95, no chroma subsampling); soft-edge mixtures vary with background. Two draws per arm yield 23,176 composites. Visible inpainted area has mean 4.13%, p95 8.45%, maximum 20.37%, and arm means 4.13%/4.12%. Correctness is averaged across donors within bird, then across birds.

Outcomes and confident flips. Table D.3 reports all groups. The confident half uses the within-slice median calibrated true-label probability among originally correct birds, including ties. In WW, opposite-type/control flip rates are 56.4%/3.3% overall and 41.8%/0.7% in that half; the fixed threshold $p \geq 0.9$ selects no WW birds.

Table D.3: Held-out interventions. Accuracies and flip rates are seed means in percent; controlled drops are same-type minus opposite-type percentage points with nominal 95% normal-approximation half-widths. Panel (b) counts are mean confident-half counts.

(a) Paired intervention arms

Group	Birds	Original	Same type	Opposite type	Controlled drop
Landbird/land	2,255	98.9	99.3	72.6	26.7 ± 1.8
Landbird/water	2,255	71.3	72.3	99.2	-26.9 ± 2.0
Waterbird/land	642	42.3	40.6	92.3	-51.7 ± 2.0
Waterbird/water	642	**92.0**	**93.2**	**40.6**	$\mathbf{52.6 \pm 1.4}$

(b) Opposite-type flips among originally correct predictions

Group	All originally correct	Confident half	Confident-half count
Landbird/land	26.8	14.8	1,122
Landbird/water	0.2	0.1	805
Waterbird/land	1.2	0.3	136
Waterbird/water	56.4	41.8	299

Matched selection comparisons. Reapplying eligibility within each seed leaves LL and WW. Minimum validation accuracy and maximum uncalibrated NLL, evaluated separately on this identical set, agree with the surviving-fraction selector in all five seeds. All-group comparisons answer a different question; landbird/water determines the model-level certificate (Table D.4).

Centering and predictive Fisher. The post hoc check fixes the checkpoints, temperatures, validation coordinates $v = (z, b)^\top$, original scales, eligible sets and intervention outcomes. No model/probe is refitted. With $\langle\cdot\rangle_g$ the divisor-n_g slice mean, compare

$$J_g = \langle vv^\top\rangle_g = \Sigma_g + \mu_g\mu_g^\top, \quad \Sigma_g = \langle (v-\mu_g)(v-\mu_g)^\top\rangle_g, \quad \mu_g = \langle v\rangle_g.$$

Centering does not rescale. The local model $\operatorname{logit} p_{\theta,\xi}(x) = f_w(x)/\widehat{T} + \theta z(x) + \xi b(x)$ at zero has score $(Y-p)v$ with conditional mean zero even when $\mu_g \neq 0$. Its predictive Fisher and the common fraction are

$$I_g = \langle p(1-p)vv^\top\rangle_g, \qquad r_\rho(A) = 1 - \frac{A_{zb}^2}{A_{zz}(A_{bb}+\rho)}, \quad \rho = 10^{-3}.$$

Perturbations follow calibration and profile only the declared background coordinate. The archived normalized-weight comparator is $K_g = I_g/\bar{q}_g$, $\bar{q}_g = \langle p(1-p)\rangle_g$, so $r_\rho(K_g) = r_{\bar{q}_g\rho}(I_g)$: normalization preserves the unregularized fraction but changes effective ridge. Table D.4 therefore reports it separately.

Table D.4: Selection and matrix sensitivity, evaluated within each seed. Panel (a) distinguishes original eligible and all-group sets; drops are mean percentage points. Panel (b) fixes the original eligible LL/WW set and interventions; fractions/drops are means ± nominal 95% normal-approximation half-widths.

(a) Matched selectors

Statistic	Candidates	Selected group, seeds	Drop
Minimum raw task information	Eligible	Landbird/land, 5/5	26.7
Minimum absolute projected information	Eligible	Landbird/land, 5/5	26.7
Minimum surviving fraction	Eligible	Waterbird/water, 5/5	52.6
Minimum validation accuracy or maximum validation NLL	Eligible	Waterbird/water, 5/5	52.6
Minimum normalized-weight fraction	Eligible	Waterbird/water, 5/5	52.6
Maximum NLL or minimum accuracy	All four	Waterbird/land, 5/5	−51.7
Group determining $I^{\mathrm{FG}}_{0.95}$	All four	Landbird/water, 5/5	−26.9

(b) Frozen-coordinate constructions

Construction	LL fraction	WW fraction	Selected, seeds	Drop
Uncentered J_g	0.590 ± 0.027	0.248 ± 0.015	WW, 5/5	52.6 ± 1.4
Centered Σ_g	0.779 ± 0.028	0.855 ± 0.032	LL, 5/5	26.7 ± 1.8
Predictive Fisher I_g	0.987 ± 0.014	0.273 ± 0.016	WW, 5/5	52.6 ± 1.4
Normalized weights K_g	0.987 ± 0.014	0.269 ± 0.016	WW, 5/5	52.6 ± 1.4

Centering selects LL in all seeds, whereas predictive Fisher retains WW. Slice means therefore contribute to the original ranking: for WW in seed 22, $\mu_z^2 = 2.094$ versus $\Sigma_{zz} = 0.0636$. Table D.5 preserves the per-seed reversal; full means, covariance entries, weights, matrices and hashes are retained in the artifact.

Table D.5: Per-seed fractions on the original global scales. Each cell is LL / WW; the smaller value determines selection. Centering reverses the ranking in every seed.

Seed	Uncentered $r_\rho(J_g)$	Centered $r_\rho(\Sigma_g)$	Predictive $r_\rho(I_g)$
22	0.5570 / 0.2291	0.7430 / 0.8542	0.9817 / 0.2587
23	0.6048 / 0.2494	0.8098 / 0.8481	0.9920 / 0.2729
24	0.6356 / 0.2755	0.8132 / 0.8359	0.9620 / 0.3032
25	0.5803 / 0.2392	0.7514 / 0.9156	0.9998 / 0.2571
26	0.5716 / 0.2460	0.7779 / 0.8189	0.9998 / 0.2732

Auxiliary flip prediction. Originally correct predictions contribute two opposite-type donor rows per bird. Baseline predictors are uncalibrated NLL, calibrated true-label confidence, absolute raw logit and standardized task logit z. Augmentation adds $\beta_g b$ and $|z - \beta_g b|$, with $\beta_g = \widehat{J}_{TN,g}/(\widehat{J}_{NN,g} +$

10^{-3}) frozen from the validation block. Both models use identical five-fold, nonshuffled CUB-species `GroupKFold` splits, keeping species and donor pairs together. Logistic regression uses an intercept, ℓ_2 penalty, $C = 1$, L-BFGS, tolerance 10^{-4}, maximum 5,000 iterations, seed 0 and no class weighting or tuning. All columns, including group indicators, are centered and scaled using training-fold means and population standard deviations; standard deviations below 10^{-9} are replaced by one.

Pool out-of-fold predictions within seed before computing average precision (AUPRC), AUROC and mean log likelihood. For clipped probabilities $q \in [10^{-12}, 1 - 10^{-12}]$,

$$G_s = \frac{1}{N_s} \sum_{i=1}^{N_s} \left[\mathrm{LL}(y_{si}, q_{si}^A) - \mathrm{LL}(y_{si}, q_{si}^B) \right], \quad \mathrm{LL}(y, q) = y \log q + (1 - y) \log(1 - q). \tag{D.2}$$

G_s is augmented-minus-baseline gain in nats per donor row. Table D.6 averages paired changes equally across seeds 22–26. The coordinates improve prediction without group indicators and within WW; with indicators, AUPRC decreases. This tests their incremental value, not unique attribution to the scalar information statistic. Original inputs, exact-source refits and paired records are in the artifact.

Table D.6: Incremental flip prediction. Counts and rates are seed means; parentheses count seeds with improvement. Changes are paired within seed.

Analysis	Rows	Flip rate	Log-likelihood gain	ΔAUPRC	ΔAUROC
All groups, no indicators	9,398	0.199	+0.016	+0.035 (5/5)	+0.096 (5/5)
All groups, with indicators	9,398	0.199	+0.001	−0.007 (0/5)	+0.003 (5/5)
Selected slice	1,181	0.564	+0.013	+0.026 (5/5)	+0.015 (5/5)

Development chronology. The surviving-fraction rule followed an absolute-information pilot on seeds 17–21 and was frozen for newly trained seeds 22–26 on the same benchmark. These tables consistently use the documented primary rerun; the original-versus-rerun ledger is retained in the artifact. Appendix E.3 gives the Fisher-Glass training comparison under this same intervention.

D.3 Synthetic Classification and Tail Checks

Draw $U_e \sim \mathrm{Uniform}(0, 1)$, set $\delta_e = 0.35 U_e^{1/\gamma}$, and draw independent standard normals s, r for each observation. The task is $z = s + \delta_e r$, $Y = \mathbf{1}\{z > 0\}$, scored by nuisance-only logit $4s$. Thus

$$\Pr(\delta_e < t) = (t/0.35)^\gamma \ (0 \le t \le 0.35), \qquad J_{TT,e} = 1 + \delta_e^2, \quad F_e^0 = \delta_e^2. \tag{D.3}$$

Population feature moments are exact; loss/accuracy use 256 observations per environment. The $\gamma = 1$ experiment uses 300,000 environments and seeds 11–15. Fit the log-survival slope over the largest 5% of inverse information; empirical CVaR averages values at or above the empirical α-quantile. Figure 2 regenerates seed 11 with $L_e = (\delta_e^2 + 10^{-12})^{-1}$, rather than recovering missing original raw output; its survival curve and the five-seed summaries have different replication levels.

Table D.7 checks the tail law and scalar independent-branch specialization of Mansouri and Arai [2026]. For K independent draws, information $K^{-1} \sum_j \delta_{e,j}^2$ has exponent $K/2$ at $\gamma = 1$; weights affect its constant. The mass sweep predicts $\gamma/2$. An anchor with $\delta_0 = 0.5$ and reserve mass 0.05 yields information $0.95\delta_e^2 + 0.05(0.25)$, floor 0.0125, and measured certificate 0.01260 ± 0.000001 (five seeds; Appendix E.1).

Table D.7: Synthetic checks. Panel (a) and the branch rows are five-seed means ± nominal 95% normal-approximation half-widths; zero widths reflect rounding. Mass-sweep rows are reported single-seed values.

(a) Prediction and information

Quantity	Estimate
Mean accuracy	$94.54 \pm 0.00\%$
Fifth-percentile environment accuracy	$88.67 \pm 0.00\%$
Loss-$\mathrm{CVaR}_{0.95}$	0.262 ± 0.000
$I^{\mathrm{FG}}_{0.95}$	$(1.22 \pm 0.53) \times 10^{-8}$
Inverse-information exponent $\widehat{\beta}$	0.498 ± 0.003

(b) Inverse-information tail exponents

Construction	Setting	Predicted β	Estimated β
Mass near the null	$\gamma = 1$	0.5	0.509
Mass near the null	$\gamma = 2$	1	0.985
Mass near the null	$\gamma = 4$	2	2.001
Independent branches	$K = 1$	0.5	0.499 ± 0.003
Independent branches	$K = 2$	1	1.000 ± 0.009
Independent branches	$K = 3$	1.5	1.508 ± 0.015

D.4 Hidden-Slice Discovery Results

Appendix C gives the authoritative blind protocol: task labels, frozen MobileNet features and a fixed CLIP direction are available; background/group metadata remain sealed through certification. Results use the corrected v2 cohort and seeds 17–21, with final full-validation gates frozen before test. Evaluation support correction may depend on the test sample. The five primary CLIP gates require no correction. Information uses same-class pooling; accuracy, uncalibrated NLL and composition use admissible within-class weights without pooling. The corrected cohort and original hashes are identified in the artifact; the older uncorrected cohort is excluded.

All selectors share classifiers with average accuracy $76.3 \pm 0.7\%$ and post hoc worst-group accuracy $57.3 \pm 1.3\%$. CLIP search finds 98.8% slice accuracy. Table D.8 includes every matched-mass comparator and paired contrast. The random-direction and both unweighted intervals contain zero; the latter reevaluate predictive-selected gates and do not rediscover them under different matrix objectives.

Class selection and overlap. CLIP, random-direction, random-slice and constant-gate selectors choose landbirds in all five seeds; high loss chooses landbirds in seeds 17,18,21 and waterbirds in 19,20. The two different-class Jaccards are zero; same-class values are 0.00222, 0.00111, and 0. Their mean 0.000666, reported as 0.001, combines between-class and within-class differences under Appendix C.1's top-count rule.

Support and sensitivity. On the common same-class training-calibration reference set, cross-fold gate-score correlation is $0.419\,[0.353, 0.485]$ and overlap $0.258\,[0.225, 0.291]$. Test ESS is $414\,[383, 445]$ using admissible weights. Raw membership mass is $0.083\,[0.077, 0.089]$ within the selected class (target 0.10), not the full test set; landbirds contain 4,510 test examples. Validation cross-fitting, separate from final testing, gives information 0.004–0.005 over masses $\{0.05, 0.10, 0.20\}$ versus high-loss 0.012–0.027, and 0.004–0.006 over shrinkages $\{0.05, 0.10, 0.20\}$.

Post-certification composition is 97.0% landbird/land for CLIP, 94.4% for random direction, and 58.4% landbird/water for high loss. CLIP information is below the weakest predefined group's in every seed. A metadata-trained oracle probe used only after blind freeze gives information 0.003 ± 0.000, NLL 0.009 ± 0.004, and 99.1% landbird/land composition (zero width is rounding). Repair outcomes are in Appendix E.3.

Table D.8: Hidden-slice comparisons. Panel (a): means ± 95% Student-t half-widths across five paired seeds; predictive information retains calibration and shrinkage $\eta = 0.10$. Panel (b): CLIP-search minus comparator, with paired 95% Student-t intervals; negative differences indicate lower information.

(a) Held-out selectors

Selector	Predictive projected information	Slice NLL
Fisher search, fixed CLIP direction	0.005 ± 0.001	0.031 ± 0.017
High loss, same target within-class mass	0.019 ± 0.006	5.779 ± 0.496
Random, same target within-class mass	0.013 ± 0.001	1.025 ± 0.059
Constant class-conditional gate	0.013 ± 0.001	1.027 ± 0.049
Fisher search, random direction	0.006 ± 0.001	0.031 ± 0.015

(b) Paired information contrasts

Information	Comparison	Difference	95% interval
Predictive	High-loss slice	-0.014	$[-0.020, -0.008]$
Predictive	Random slice	-0.008	$[-0.009, -0.007]$
Predictive	Constant gate	-0.008	$[-0.009, -0.007]$
Predictive	Random direction	-0.0004	$[-0.0015, 0.0007]$
Unweighted feature	High-loss slice	-0.034	$[-0.106, 0.039]$
Centered unweighted	High-loss slice	-0.071	$[-0.166, 0.024]$

D.5 Predefined-Group Diagnostics

Fitting. CivilComments-WILDS [Borkan et al., 2019, Koh et al., 2021] uses training-fitted unigram/bigram TF-IDF (50,000 features, minimum document frequency 5, sublinear TF) and class-balanced SGD logistic task/identity-any probes ($\ell_2 = 10^{-5}$, 50 passes, tolerance 10^{-3}). Identity-by-label groups require 100 examples; tail regression uses the largest 50% of group inverse information, requiring eight groups and four tail points. Seed 13 gives the fixed diagnostic; seeds 13–17 give aggregates.

Waterbirds uses frozen ImageNet MobileNetV3-small features (extraction batch 64) and balanced L-BFGS logistic task/background models (1,000 maximum iterations). Four groups are reported directly, without fitting a tail exponent; only controls/anchors vary over seeds 17–21. Primary task/probe and metadata coordinates use training moments. These original runs used an i9-11900F/32 GB workstation, about one minute for CivilComments and three minutes for first Waterbirds extraction/fitting, then cached features.

Table D.9: Original predefined-group diagnostics, distinct from fresh audits in Appendix B.4. Means ± nominal 95% normal-approximation half-widths; Waterbirds zeros reflect fixed-model fitting.

Dataset	Average accuracy (%)	Worst-group accuracy (%)	Loss-$\mathrm{CVaR}_{0.95}$	$I^{\mathrm{FG}}_{0.95}$
CivilComments	86.10 ± 0.22	43.75 ± 0.64	0.998 ± 0.012	0.500 ± 0.001
Waterbirds	77.11 ± 0.00	56.54 ± 0.00	1.741 ± 0.000	0.165 ± 0.000

For fixed CivilComments, the weakest-information group is `muslim`/$y = 0$ (accuracy 60.1%); lowest accuracy is `black`/$y = 0$ (44.0%). Waterbirds' weakest-information group is landbird/land (97.3% accuracy), whereas waterbird/land has lowest accuracy (56.5%); inverse-information p95 is 5.90.

Specification and stress controls. The prespecified semantic family is $\{b_{\mathrm{meta}}, b_{\mathrm{probe}}\}$: standardized identity-any/background annotation and its training-fitted probe, without cross-fitting. Semantic relevance follows those annotations; no held-out relevance threshold, anchor task-preservation tolerance, fold-stability admission screen, or nuisance-conditioning threshold was applied. The ridge is 10^{-3}; family minima are 0.499 and 0.165. Reserve anchors (identity-present label-balanced

CivilComments; four-group-balanced Waterbirds) are construction inputs, not admission samples. Fresh exact/ridge audits and hidden-gate stability are separate analyses.

Table D.10: Nuisance sensitivity of fixed-model diagnostics. Entries are $I^{\mathrm{FG}}_{0.95}$; worst-group accuracy remains 44.0%/56.5% for CivilComments/Waterbirds.

Specification	CivilComments	Waterbirds	Role
Metadata	0.587	0.266	Semantic-family member
Learned probe	0.499	0.165	Semantic-family minimum
Random	0.588	0.482	Negative control
Under-projected	0.588	0.482	Under-projection control
Over-projected	0.001	0.012	Task-erasure control
Hidden-loss	0.060	0.177	Loss-coordinate stress control

Random nuisance is independent standard normal; under-projection sets it to zero. Over-projection uses standardized task score plus Gaussian noise of standard deviation 0.01 and deliberately erases task information. "Hidden-loss" is observed test loss with original groups fixed. The random, over-projected, and observed-loss coordinates use evaluated-array normalization and are not independently learned nuisance scores. The metadata/random contrast is 0.587/0.588 on CivilComments, versus probe/random 0.165/0.482 on Waterbirds, illustrating specification dependence.

The versioned artifact (see the artifact availability statement) retains execution commands, full-precision matrices and overlaps, run-history ledgers and expanded source tables, with original cohort provenance. Those records supplement the scientific definitions and controls above.

E Information-Changing Interventions

This appendix gives the reserve proof, training and acquisition comparisons, and the influence derivative behind Section 6. Distinct cohorts and information targets remain separate. Expanded protocols and numerical breakdowns are retained in the versioned artifact's `docs/extended_appendix_details/`, with executable reproduction guides and provenance (see the artifact availability statement).

E.1 Information Reserves under Observed-Source Allocation

Fix the predictor, coordinates, and normalization. Let $J_e, J_0 \succeq 0$ have positive-definite nuisance blocks, and suppose the anchor satisfies $F_0^0 := S_0(J_0) \succeq \rho_0 I$, $\rho_0 > 0$. An observed-source allocation instantiates the Fisher-reserve principle of Mansouri and Arai [2026]. The source label is observed and allocation probabilities are parameter-independent. Sampling the anchor with probability $r \in (0, 1]$ therefore gives

$$\begin{aligned} J_e^{\mathrm{res}} &= (1-r)J_e + rJ_0, \\ S_0(J_e^{\mathrm{res}}) &\succeq (1-r)S_0(J_e) + rS_0(J_0) \succeq r\rho_0 I. \end{aligned} \tag{E.1}$$

The same statement applies to affine allocation of a fixed operational matrix map; it is not an assertion about an unobserved mixture likelihood.

Proof. For $J = \left(\begin{smallmatrix} A & B \\ B^\top & C \end{smallmatrix}\right)$, $C \succ 0$, completing the square gives

$$u^\top S_0(J)u = \min_v q_J(u, v), \qquad q_J(u, v) = \begin{pmatrix} u \\ -v \end{pmatrix}^\top J \begin{pmatrix} u \\ -v \end{pmatrix}, \quad v_* = C^{-1}B^\top u. \tag{E.2}$$

Linearity in J implies

$$\begin{aligned} \min_v\{(1-r)q_{J_e}(u, v) + rq_{J_0}(u, v)\} &\geq (1-r)\min_v q_{J_e}(u, v) + r\min_v q_{J_0}(u, v) \\ &\geq r\rho_0\|u\|^2. \end{aligned} \tag{E.3}$$

The last inequality uses $S_0(J_e) \succeq 0$ and the anchor bound; holding for every u, it proves (E.1). □

For any environment distribution, $0 < \alpha < 1$ and $\varepsilon \geq 0$, CVaR monotonicity gives

$$L_e^{\mathrm{res},(\varepsilon)} \leq (r\rho_0 + \varepsilon)^{-1}, \qquad I_\alpha^{\mathrm{FG,res}}(\varepsilon) \geq r\rho_0 + \varepsilon \geq r\rho_0. \tag{E.4}$$

Here $r = m/(n+m)$ for m anchor and n base observations. An estimated anchor's positive lower bound from Appendix B.3 yields the same guarantee on its confidence event. The claim concerns the fixed predictor; retraining requires a new audit. Appendix D.3 evaluates $r = .05$, $\rho_0 = .25$, and the resulting floor .0125.

E.2 Regularized Training

For a declared group matrix $\widehat{J}_g(w)$, put $\widehat{L}_g(w) = [\lambda_{\min}(S_{\mathrm{rid}}(\widehat{J}_g(w))) + \varepsilon]^{-1}$. With fixed $\pi_g \geq 0$, $\sum_g \pi_g = 1$, training solves

$$\min_{w,\tau} \widehat{R}_{\mathrm{base}}(w) + \lambda_{\mathrm{FG}}\left[\tau + \frac{1}{1-\alpha}\sum_g \pi_g(\widehat{L}_g(w) - \tau)_+\right], \tag{E.5}$$

where the base risk is ERM or Group DRO. Frozen-feature Waterbirds uses unweighted second moments of the standardized task logit and fixed background-probe score, ridge 10^{-3} and $\varepsilon = 10^{-8}$, rather than the hidden-slice predictive-$p(1-p)$ matrix. Training $\alpha = .75$ and audit $\alpha = .95$ both give the maximum inverse information over four equally weighted groups.

Protocol and selection. ImageNet MobileNetV3-small features and background-probe scores use training-moment standardization. A linear head receives 1,200 full-batch AdamW updates (lr $= .01$, weight decay 10^{-4}, seeds 17–21), Gaussian $N(0, .01^2)$ initial weights and zero intercept. Current logits use full-training population standard deviation plus 10^{-6}; training projected information is floored at ε before reciprocal stabilization. Evaluation uses fitted training-logit moments. Group DRO minimizes the maximum group mean cross-entropy; Loss-CVaR uses the largest $\lceil .1n \rceil$ training losses, and JTT refits with weight 20 on initial ERM mistakes.

For each FG family and seed, search $\Lambda_{\mathrm{FG}} = \{.01, .03, .1, .3\}$. Among candidates whose validation worst-group accuracy is within .02 of the family's best, maximize the validation certificate, breaking ties by higher validation worst-group accuracy. This uses validation group annotations. Historical choices are .03 for standalone FG in all seeds and .1 or .3 for Group DRO+FG. Table E.1 retains all six methods; official test data enter only final evaluation.

Table E.1: Original frozen-feature training cohort, seeds 17–21. Both panels give means ± nominal 95% normal-approximation half-widths; accuracy is percent and the ridge certificate uses $\alpha = .95$.

Method	Avg. acc. (%)	Worst acc. (%)	Loss-CVaR$_{.95}$
ERM	77.2 ± 0.1	55.7 ± 0.3	1.969 ± 0.007
Group DRO	80.9 ± 0.3	67.3 ± 0.3	1.143 ± 0.021
JTT	77.4 ± 0.1	54.1 ± 0.2	2.167 ± 0.004
Loss-CVaR	78.2 ± 0.1	52.1 ± 0.3	2.427 ± 0.041
Fisher-Glass	80.4 ± 0.2	56.3 ± 0.3	1.866 ± 0.025
Group DRO + FG	80.6 ± 0.9	69.4 ± 1.3	0.855 ± 0.082
Method	**Ridge certificate**	**Inverse p95**	
ERM	0.198 ± 0.001	4.86 ± 0.02	
Group DRO	0.233 ± 0.003	4.10 ± 0.06	
JTT	0.234 ± 0.001	4.13 ± 0.01	
Loss-CVaR	0.233 ± 0.002	4.17 ± 0.03	
Fisher-Glass	0.399 ± 0.006	2.46 ± 0.04	
Group DRO + FG	0.614 ± 0.059	1.63 ± 0.17	

FG doubles ERM's certificate; adding it to Group DRO raises the certificate from .233 to .614 and mean worst-group accuracy from 67.3% to 69.4%. These historical summaries differ from Appendix B's fresh ridge/exact audit cohort, which does not retroactively certify the historical checkpoints. Their original per-seed outputs were not recovered; retained reports, complete settings and independent rerun commands are in the artifact (see the artifact availability statement).

E.3 Extended Training and Selection Results

End-to-end penalty frontier. MobileNetV3-small is fine-tuned for 10 epochs (batch 128, evaluation batch 512, seeds 17–21) with AdamW learning rate and weight decay 10^{-4}; the final epoch is evaluated. The reported H100 configuration uses bf16 and TF32. The nuisance is training-standardized background annotation; task logits use minibatch population standard deviation plus 10^{-6}. Groups with fewer than two observations are omitted. Ridge is 10^{-3}, $\alpha = .95$, and projected information is floored at 10^{-8} before another 10^{-8} enters the reciprocal. Table E.2 retains the complete frontier: stronger penalties can increase information while degrading prediction.

At 10^{-4} the certificate rises from .259 to .502 and inverse p95 falls from 3.68 to 2.10; at 10^{-2} worst-group accuracy is only 22.7%. Historical report summaries support this table; Appendix B's retained fresh audits are a separate cohort.

Prospective information training. Use Appendix D.2's shared model, 50-epoch schedule and seeds 22–26. Backbone fine-tuning uses AdamW with learning rate and weight decay 10^{-4}. FG adds coefficients $\{.003, .01, .03\}$ and a background probe trained on detached features (AdamW learning rate .001, default weight decay .01), whose scores are detached in the penalty. After one penalty-free epoch, training uses $\alpha = .75$, ridge .001, batch standard-deviation guard 10^{-6}, groups of at least six observations with nuisance standard deviation at least .05, and a .05 information floor before direct inversion. This floor differs from audit stabilization.

Table E.2: End-to-end Waterbirds frontier. Each non-reference row uses Group DRO + FG. FG entries are means ± sample standard deviations over five seeds, not confidence intervals. The Group DRO reference is the reported comparator at $\lambda_{\mathrm{FG}} = 10^{-4}$; each sweep reran its baseline, so this is not a shared paired reference for every coefficient.

Objective / λ_{FG}	Avg. acc. (%)	Worst acc. (%)	Loss-CVaR$_{.95}$
Group DRO reference	85.4	53.2	0.925
5×10^{-5}	79.7 ± 6.2	56.6 ± 5.8	0.733 ± 0.061
10^{-4}	79.5 ± 10.2	57.7 ± 8.3	0.733 ± 0.110
3×10^{-4}	69.9 ± 9.7	58.8 ± 8.7	0.653 ± 0.064
10^{-3}	62.7 ± 7.9	49.9 ± 1.8	0.710 ± 0.012
3×10^{-3}	57.7 ± 4.5	48.6 ± 6.6	0.717 ± 0.036
10^{-2}	47.9 ± 16.5	22.7 ± 16.7	1.031 ± 0.416
Objective / λ_{FG}	Ridge certificate	Inverse p95	
Group DRO reference	0.259	3.68	
5×10^{-5}	0.393 ± 0.059	2.49 ± 0.34	
10^{-4}	0.502 ± 0.166	2.10 ± 0.58	
3×10^{-4}	0.614 ± 0.189	1.71 ± 0.46	
10^{-3}	0.786 ± 0.124	1.29 ± 0.20	
3×10^{-3}	0.817 ± 0.068	1.22 ± 0.10	
10^{-2}	0.809 ± 0.076	1.23 ± 0.12	

Select the final model with largest validation certificate among candidates whose overall validation accuracy is at least ERM minus .01; an empty set triggers a recorded unconstrained fallback. Choices are $[.01, .03, .01, .003, .01]$, with fallbacks in seeds 22 and 24. All audit objects are frozen before intervention scoring. In the current v2 rerun, FG raises the validation certificate from .289 to .330, but the controlled drop is 52.9 versus 52.6 percentage points for ERM and the originally-correct flip rate is 57.7% versus 56.4%. Thus accuracy matching fails in two seeds and background sensitivity persists.

Constrained hidden-slice repair. The corrected cohort of Appendix D.4 compares loss, Fisher and combined repair without opening group metadata. Each arm includes the exact unchanged model and must preserve overall, balanced and each-class validation accuracy within .02 of the base. Loss coefficients are $\{.1, .3, 1\}$, Fisher coefficients $\{.01, .03, .1\}$, and combined candidates use their Cartesian product. The same warm start and fresh AdamW optimizer (lr $= .01$, weight decay 10^{-4}, gradient clip 10) receive at most 1,200 steps. Calibration NLL at temperature one is checked every 25 steps, with patience 12, improvement threshold 10^{-5}, and step-zero-eligible best-checkpoint restoration. The matched ERM continuation sets both penalties to zero. Base temperature remains fixed for information training, selection and subsequent audits.

For the frozen training slice, the dimensionless Fisher penalty is

$$P_{\mathrm{FG}}(w) = \frac{\max\{F_{\mathrm{init}}, \varepsilon\}}{F(w) + \varepsilon}, \qquad \varepsilon = 10^{-8}, \tag{E.6}$$

with predictive shrinkage .1, ridge .001, and negative numerical information clipped to zero. Within the accuracy envelope, loss/Fisher repair minimizes frozen-slice loss/maximizes information. Combined selection uses the loss–information Pareto frontier, maximizing the smaller relative improvement; ties favor mean relative improvement, smaller normalized coefficients, then earlier index. Fresh rediscovery after selection uses matched gate/fold seeds, so the reported population can change.

Fisher repair chooses the unchanged model in every seed. Paired worst-group changes versus matched continuation are $-.4$, .0 and $-.6$ points for loss, Fisher and combined repair, with intervals containing zero. The oracle reaches 64.8% worst-group accuracy while rediscovered information remains .005. LL denotes landbird/land.

Post hoc group-free coefficient selection. This analysis was designed after earlier group/test outcomes were inspected. Its grid $\{0, .01, .03, .1, .3\}$ uses (E.6), not the known-group objective's scale. Training is partitioned $70/15/15\%$ for fitting/calibration/gate discovery; a frozen train-only

Table E.3: Corrected constrained repair, seeds 17–21. Accuracy, NLL and information are means $\pm$ 95% Student-t half-widths; composition is a mean percentage. Panel two audits each arm's rediscovered CLIP-defined slice. The oracle uses group annotations.

Arm	Avg. acc. (%)	Worst acc. (%)	Unchanged
Matched ERM continuation	76.3 ± 0.7	57.3 ± 1.3	—
Loss repair	76.3 ± 0.6	56.9 ± 1.3	$3/5$
Fisher repair	76.3 ± 0.7	57.3 ± 1.3	$5/5$
Combined repair	76.4 ± 0.9	56.7 ± 1.6	$4/5$
Oracle Group DRO	80.0 ± 0.3	64.8 ± 0.9	—
Arm	**Slice NLL**	**Projected info.**	**LL mass (%)**
Matched ERM continuation	0.031 ± 0.017	0.005 ± 0.001	97.0
Loss repair	0.023 ± 0.017	0.005 ± 0.002	97.8
Fisher repair	0.031 ± 0.017	0.005 ± 0.001	97.0
Combined repair	0.031 ± 0.017	0.005 ± 0.001	97.2
Oracle Group DRO	0.033 ± 0.018	0.005 ± 0.002	95.9

slice supplies the penalty. All candidates share the warm start, training seed, optimizer, 1,200-step budget, checkpoint rule and base temperature above. Official validation enters only after heads are fixed; group metadata stays sealed through blind test evaluation.

Let $\widehat{R}_\lambda$ be mean fixed-temperature validation NLL, s_λ its per-example standard error, and $\lambda_R = \arg\min_\lambda(\widehat{R}_\lambda, \lambda)$ lexicographically. Five task-stratified validation folds rediscover candidate-specific slices; $V_{k,\lambda}$ is held-out inverse predictive information and $s_{V,\lambda}$ its sample standard deviation. Selection uses

$$\begin{aligned} \Lambda_{\text{risk}} &= \{\lambda : \widehat{R}_\lambda \le \widehat{R}_{\lambda_R} + s_{\lambda_R}\}, \quad U^{\text{adj}}_\lambda = \overline{V}_\lambda + t_{4,.99} s_{V,\lambda}/\sqrt{5}, \\ \widehat{\lambda}_{\text{hidden}} &= \min \arg \max_{\lambda \in \Lambda_{\text{risk}}} (U^{\text{adj}}_\lambda)^{-1}, \quad t_{4,.99} = 3.74695. \end{aligned} \tag{E.7}$$

The risk allowance uses the minimizer's SE for all candidates. Bonferroni adjustment covers five coefficients, but overlapping gate-training folds make this a ranking score without independent-fold coverage. Accuracy screens and paired inverse-information tests are diagnostics, not selection constraints.

Table E.4: Post hoc coefficient selection, five paired seeds. Accuracies are test results; hidden information is predictive projected information on test slices defined by the corresponding validation-rediscovered, frozen gates. Brackets are 95% Student-t intervals. Policies selecting identical models across these five seeds share rows; the text explains both coincidences. Coefficient counts are λ : number of seeds.

Policy	Average accuracy (%)	Worst accuracy (%)
Frozen ERM / matched zero	75.5 [74.3, 76.6]	55.7 [53.3, 58.2]
Fixed $\lambda = .03$	76.6 [76.1, 77.2]	51.5 [47.1, 55.8]
Risk only / group-free	77.5 [76.5, 78.6]	43.1 [33.1, 53.1]
Group-supervised comparator	77.7 [76.6, 78.8]	44.2 [34.2, 54.3]
Policy	**Coefficient counts**	**Hidden information**
Frozen ERM / matched zero	— / $0:5$	0.007 [0.005, 0.009]
Fixed $\lambda = .03$	$.03:5$	0.007 [0.004, 0.010]
Risk only / group-free	$0:1,\ .3:4$	0.012 [0.006, 0.019]
Group-supervised comparator	$0:1,\ .1:1,\ .3:3$	0.011 [0.005, 0.018]

Matched zero continuation restores step zero in all five seeds, explaining equality with frozen ERM despite intervening optimization. Group-free selection equals plain NLL minimization in every seed; the information criterion does not change any choice. Versus matched continuation, paired mean changes are $+2.1$ average-accuracy points ($[+.2, +4.0]$), $+.006$ information ($[+.001, +.010]$), and -12.6 worst-group points ($[-21.6, -3.7]$). The group-supervised comparator opens validation annotations only after blind decisions/artifacts are fixed and is not an upper bound. These schema-

v2 results and the corrected repair cohort are separately identified in the artifact (see the artifact availability statement); extended settings and full tables are retained there.

E.4 Tail-Transverse Influence

Fix the predictor, coordinates, calibration, slice rules and finite audit weights $\pi_g \geq 0$, $\sum_g \pi_g = 1$. An action specifies a positive-semidefinite path $J_g(t,a) = J_g + tH_{g,a} + o(t)$ as $t \downarrow 0$, with symmetric $H_{g,a}$ and an operator-norm remainder. Allocation, normalization and retained pooling belong in this path; audit weights stay fixed.

One differential for exact and ridge targets. Write $J_g = \left(\begin{smallmatrix} A_g & B_g \\ B_g^\top & C_g \end{smallmatrix}\right)$. Choose a fixed $\rho \geq 0$ with $C_g + \rho I \succ 0$: $\rho = 0$ gives exact Schur information and $\rho = \lambda_{\mathrm{rid}}$ gives the ridge target. Set $Q_g = (C_g + \rho I)^{-1}$, $S_\rho(J_g) = A_g - B_g Q_g B_g^\top$, and let λ_g be its smallest eigenvalue. Since $DQ_g[H_{NN}] = -Q_g H_{NN} Q_g$,

$$\begin{aligned} DS_\rho(J_g)[H] &= H_{TT} - H_{TN}Q_gB_g^\top - B_gQ_gH_{TN}^\top \\ &\quad + B_gQ_gH_{NN}Q_gB_g^\top, \\ p_g(u) &= \begin{pmatrix} u \\ -Q_gB_g^\top u \end{pmatrix}, \qquad u^\top DS_\rho(J_g)[H]u = p_g(u)^\top H p_g(u). \end{aligned} \tag{E.8}$$

For a simple minimum eigenvalue with unit eigenvector u_g, differentiating the eigenvalue equation therefore gives

$$h_{g,a} := D_a\lambda_g = p_g^\top H_{g,a} p_g, \qquad p_g = p_g(u_g). \tag{E.9}$$

For allocation rate $\eta_g \geq 0$, additive and normalized contributions have different derivatives:

$$\begin{aligned} H_{g,a} = \eta_g s_a s_a^\top &\implies h_{g,a} = \eta_g (p_g^\top s_a)^2, \\ J_g(t,a) = \frac{J_g + t\eta_g M_{g,a}}{1 + t\eta_g} &\implies H_{g,a} = \eta_g (M_{g,a} - J_g), \\ &\qquad h_{g,a} = \eta_g \left[p_g^\top M_{g,a} p_g - \lambda_g + \rho \|Q_g B_g^\top u_g\|^2 \right]. \end{aligned} \tag{E.10}$$

Indeed $Q_g C_g Q_g = Q_g - \rho Q_g^2$, so $p_g^\top J_g p_g = \lambda_g - \rho\|Q_g B_g^\top u_g\|^2$; the correction vanishes at $\rho = 0$. The additive rank-one gain is η_g times the squared residual $[u_g^\top (s_{T,a} - B_g Q_g s_{N,a})]^2$. This is the efficient residual for $\rho = 0$ and the corresponding ridge residual for $\rho > 0$. Replacement instead uses the allocated new-minus-removed contribution. For one observation appended to group g, its indicator and original sample-count scaling enter η_g; they cannot be omitted when claiming a normalized-augmentation derivative.

Tail sensitivity, including nonsmooth cases. Fix $\varepsilon > 0$ and $0 < \alpha < 1$, and put $L_g = (\lambda_g + \varepsilon)^{-1}$. Appendix A.4's finite CVaR dual uses $\mathcal{W}_{\alpha,\pi} = \{\omega \geq 0 : \sum_g \omega_g = 1,\ \omega_g \leq \pi_g/(1-\alpha)\}$. Let $\Omega_\alpha^\star(L)$ maximize $\sum_g \omega_g L_g$ over this set. If the maximizer ω_α is unique and eigenvalues are simple,

$$G_\alpha(a) := -D_a \operatorname{CVaR}_{\alpha,\pi}(L) = \sum_g \omega_{\alpha,g} \frac{h_{g,a}}{(\lambda_g + \varepsilon)^2}, \qquad D_a I_{\alpha,\pi} = I_{\alpha,\pi}^2 G_\alpha(a). \tag{E.11}$$

This follows from $D_a L_g = -h_{g,a}/(\lambda_g + \varepsilon)^2$ and differentiation of the unique active linear functional. The dual weights include fractional threshold mass and the upper-tail normalization.

For a repeated minimum eigenvalue with eigenspace $\mathcal{E}_g$, symmetric eigenvalue perturbation gives the first formula below. At a CVaR tie, the right derivative of a maximum is the maximum of its active directional derivatives; negation gives the second:

$$\begin{aligned} h_{g,a} &= \min_{\substack{u \in \mathcal{E}_g \\ \|u\|=1}} p_g(u)^\top H_{g,a} p_g(u), \\ G_\alpha^+(a) := -D_a^+ \operatorname{CVaR}_{\alpha,\pi}(L) &= \min_{\omega \in \Omega_\alpha^\star(L)} \sum_g \omega_g \frac{h_{g,a}}{(\lambda_g + \varepsilon)^2}. \end{aligned} \tag{E.12}$$

These formulas cover either or both sources of nondifferentiability; the reciprocal-certificate identity holds with right derivatives. An arbitrary fixed maximizing tail vector is only a score unless it attains this direction-dependent minimum.

For assigned label y_a, tail-transverse influence multiplies the directional information gain by current prediction loss:

$$\mathrm{TTI}(a) = G_\alpha(a)\ell_w(x_a^{\mathrm{cf}}, y_a), \qquad \ell_w(x,y) = -\log p_w(y \mid x), \tag{E.13}$$

using G_α^+ at nonsmooth points. This ranks actions at the fixed predictor; finite augmentation, refitting and changed normalization require a new audit. The implemented residual-square surrogate $\widetilde{G}_\alpha$ below omits allocation/baseline terms and is distinguished from this derivative.

E.5 Feature-Space Acquisition Experiments

Candidates, budgets and refitting. Using training-standardized frozen MobileNetV3-small representations, let $\bar{v}_{y,b}$ denote the official-training mean for task label y and background $b \in \{0,1\}$. Each official-validation candidate becomes

$$v^{\mathrm{cf}} = v + \gamma_{\mathrm{cf}}(\bar{v}_{y,1-b} - \bar{v}_{y,b}), \qquad y^{\mathrm{cf}} = y, \quad b^{\mathrm{cf}} = 1 - b. \tag{E.14}$$

This preserves the assigned label, not necessarily semantic or pixel-level content. The reported unit-strength run uses label-conditional means for every candidate. A balanced logistic background probe, fitted and standardized on training data, stays fixed and scores the translated features. Test data supply final evaluation only.

Budgets $b_{\mathrm{acq}} = .01, .02, .05$ use $k = \min\{\max(1, \mathrm{round}(b_{\mathrm{acq}} n_{\mathrm{train}})), n_{\mathrm{cand}}\}$, giving $48, 96, 240$ candidates from the 1,199-example validation pool relative to 4,795 original training examples. These are formula-implied counts; historical selected-index records were not recovered. Acquired features and labels augment training, and a balanced logistic task head is fitted from scratch, without an FG penalty or special acquisition weights. Both probe and head use L-BFGS, at most 1,000 iterations, and library defaults ($C = 1$, tolerance 10^{-4}, fitted intercept). Class weights are recomputed on the augmented set. New task logits are normalized by augmented-training moments, while nuisance normalization stays fixed. Acquired validation examples are no longer an independent evaluation sample. Seeds are 17–21; the historical scikit-learn version was not pinned.

Implemented score and its scope. Let (z, q) be the task logit and nuisance score standardized on original training data. Original validation-group moments are unweighted second moments, without $p(1-p)$; put $\beta_g = J_{TN,g}/(J_{NN,g} + .001)$, $\lambda_g = J_{TT,g} - J_{TN,g}^2/(J_{NN,g} + .001)$, and $L_g = (\lambda_g + \varepsilon)^{-1}$ with $\varepsilon = 10^{-8}$. Numerical negative information is floored at zero. For the threshold-inclusive tail $T = \{g : L_g \geq \mathrm{quantile}_{.95}(L_1, \ldots, L_4)\}$, the code ranks by

$$\widetilde{G}_\alpha(a) = \frac{\sum_{g \in T}(\lambda_g + \varepsilon)^{-2}(z_a^{\mathrm{cf}} - \beta_g q_a^{\mathrm{cf}})^2}{\max\{\sum_{g \in T}(\lambda_g + \varepsilon)^{-2}, 10^{-12}\}}, \qquad \widetilde{\mathrm{TTI}}_\alpha(a) = \widetilde{G}_\alpha(a)\ell_w(v_a^{\mathrm{cf}}, y_a). \tag{E.15}$$

Every tail projector is applied to every candidate, irrespective of its assigned group; for a unique worst group the inverse-square weight cancels. The score lacks group allocation, sample-count scaling, old-information subtraction and replacement terms. Its nonnegative residual square is an unnormalized rank-one contribution at fixed coordinates, not the derivative of normalized augmentation or of the post-refit certificate. Threshold-inclusive ties also differ from the fractional/direction-dependent CVaR convention in Appendix E.4.

High loss and pure tail use ℓ_w and $\widetilde{G}_\alpha$. Rank product multiplies ordinal rank-normalizations of these two scores; uncertainty product multiplies $\widetilde{G}_\alpha$ by $p(1-p)$. Equal values are not assigned averaged ranks. The acquisition baseline has 78.2% average and 52.5% worst-group accuracy, certificate .209, and weakest-information group landbird/land; it is distinct from the predefined-diagnostic cohort.

At $\gamma_{\mathrm{cf}} = 1$, the loss product exceeds high loss in mean worst-group accuracy at all budgets, while high loss gives larger certificate gains. This advantage is strength-dependent: high loss wins accuracy at every budget for .25 and 1.5, and at the 5% budget for .5. The rank product also exceeds the chosen loss product at 1% (5.45 versus 4.52 points); the chosen product is not uniformly best. Information gains and predictive gains are distinct outcomes.

Table E.5: Feature-space acquisition, means over seeds 17–21. Δa is worst-group accuracy gain in percentage points and ΔI the absolute ridge-certificate gain. Panel (a) gives the complete strength-by-budget comparison; panel (b) retains pure-tail and alternative-product controls at unit strength. Only aggregate reports survive, so historical seed-level intervals are unavailable.

(a) Strength and budget: implemented geometric score × loss versus high loss

		$\widetilde{\mathrm{TTI}}_\alpha$		High loss	
γ_{cf}	Budget	Δa	ΔI	Δa	ΔI
0.25	1%	+4.05	+0.0308	+5.76	+0.0411
	2%	+7.17	+0.0366	+7.63	+0.0477
	5%	+8.88	+0.0521	+10.12	+0.0590
0.5	1%	+6.39	+0.0205	+4.36	+0.0328
	2%	+5.61	+0.0300	+4.83	+0.0446
	5%	+6.54	+0.0417	+8.57	+0.0437
1	1%	+4.52	+0.0317	+3.27	+0.0603
	2%	+6.85	+0.0452	+5.76	+0.0581
	5%	+7.94	+0.0606	+7.01	+0.0630
1.5	1%	+7.48	+0.0662	+7.63	+0.1018
	2%	+8.72	+0.0826	+9.81	+0.1006
	5%	+9.97	+0.1159	+12.31	+0.1137

(b) Unit-strength controls

	Pure tail		Rank product		Score $\times p(1-p)$	
Budget	Δa	ΔI	Δa	ΔI	Δa	ΔI
1%	+0.62	+0.0045	+5.45	+0.0287	+1.56	+0.0012
2%	+2.80	+0.0098	+6.54	+0.0326	+1.71	+0.0073
5%	+3.58	+0.0340	+6.07	+0.0516	+3.12	+0.0230

The retained acquisition reports support Table E.5; its values are not a fresh rerun. Per-seed JSON/CSV, selected indices and pinned historical versions remain unavailable, so no historical uncertainty is reconstructed. Independent rerun code, aggregate selected-candidate score/loss summaries, allocation mechanics and expanded protocol are available in the artifact (see the artifact availability statement). Five seeds on fixed splits do not constitute five independent datasets.

Declarations and Artifact Availability

Funding. This work received no funding.

Competing interests. The author declares no competing interests.

Code and supporting artifacts. The public repository and versioned release `camera-ready-v5` are available at

https://github.com/E-zClap/fisher-glass-learning/releases/tag/camera-ready-v5.

The release provides experiment code, dependencies, derived numerical inputs, per-seed records, checksums, and commands for reproducing the retained analyses and independently rerunning the experiments. Its reproduction guides specify data preparation, environments, and experiment settings. Its artifact manifest maps manuscript tables to the supporting files and identifies the corrected hidden-slice cohort. Expanded appendix sources retain displaced execution details and numerical records; the manuscript map records consolidated table labels. The prospective matrix and auxiliary flip-prediction reanalyses can be replayed from the included frozen numerical inputs. Original benchmark images and pretrained weights are not redistributed; dataset sources and access terms are documented separately. Historical synthetic and frozen-training summaries and the acquisition results retain their stated provenance: where original per-seed outputs are unavailable, the release includes the supporting aggregate reports and executable rerun workflows; fresh experiment replication does not imply exact replay of missing historical outputs.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: Sections 1–5 distinguish conditional classification bounds, operational information targets, and empirical sensitivity findings. Section 8 states the limits on generalization, nuisance attribution, and behavioral repair.

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Section 8 discusses specification dependence, local-theory assumptions, nominal audit coverage, development history, benchmark scope, and the distinction between information gains and predictive improvement.

   Guidelines:

   - The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.

- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [Yes]

   Justification: The theoretical statements give their assumptions, with proofs in Appendices A, B, C, and E. The classification result explicitly retains distinguishability, decision-separation, and local-radius conditions.

   Guidelines:

   - The answer [N/A] means that the paper does not include theoretical results.
   - All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
   - All assumptions should be clearly stated or referenced in the statement of any theorems.
   - The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
   - Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
   - Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: [Yes]

   Justification: Appendices B–E specify the data splits, models, objectives, hyperparameters, seed sets, selection rules, interventions, and evaluation procedures. The public release provides training and evaluation code, experiment-specific rerun instructions, and numerical replay of the retained main-result artifacts. Independent reruns are distinguished from exact replay of historical outputs that were not retained.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: [Yes]

   Justification: The public `camera-ready-v5` release provides the experiment code, dependencies, data-access and preparation instructions, commands for training and evaluation, and retained numerical inputs and outputs. Its reproduction guides map the main results to these workflows and identify historical outputs unavailable for exact replay. Public benchmark data and pretrained weights are obtained from their original providers.

   Guidelines:

   - The answer [N/A] means that paper does not include experiments requiring code.
   - Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
   - While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
   - The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
   - The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.

- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: [Yes]

   Justification: Appendices C–E give the data roles, fitting and calibration procedures, optimizers, hyperparameters, seed sets, and selection rules. They distinguish original cohorts, fresh audit replications, and post hoc reanalyses.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: [Yes]

   Justification: Appendix D.1 defines nominal normal-approximation seed intervals and distinguishes Student-t intervals, reported standard deviations, and conditional bootstrap uncertainty; Appendix B.4 separately qualifies nominal Wald values. Acquisition tables explicitly report means without intervals because their per-seed outputs are unavailable.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
   - The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
   - The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
   - The assumptions made should be given (e.g., Normally distributed errors).
   - It should be clear whether the error bar is the standard deviation or the standard error of the mean.
   - It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
   - For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
   - If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [No]

   Justification: Appendices D, E.2, and E.3 identify the reported CPU/GPU hardware and available memory and timing information. Per-run execution times, peak memory, and total compute including development runs were not recorded comprehensively.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
   - The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics https://neurips.cc/public/EthicsGuidelines?

   Answer: [Yes]

   Justification: The work uses synthetic data and existing research benchmarks, with no participant recruitment or release of benchmark images or pretrained weights. Section 8 and the experimental appendices disclose the scope of the evidence and distinguish approximate numerical audits from formal guarantees.

   Guidelines:

   - The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
   - If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
   - The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [Yes]

    Justification: The audits may help identify subpopulations vulnerable to nuisance changes, but treating a restricted audit as a universal guarantee could create false assurance and obscure harms in unaudited populations. Section 8 limits conclusions to declared populations and specifications and calls for prospective validation of deployment objectives.

    Guidelines:

    - The answer [N/A] means that there is no societal impact of the work performed.
    - If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
    - Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
    - The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to

generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

    Answer: [N/A]

    Justification: The work does not release models or datasets with a high risk of misuse. The accompanying research code and numerical artifacts exclude benchmark images and pretrained model weights.

    Guidelines:

    - The answer [N/A] means that the paper poses no such risks.
    - Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
    - Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
    - We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [No]

    Justification: The original datasets, models, and methods are credited, and `DATA_LICENSES.md` links upstream terms, but complete license and version records for every historical dependency were not retained. The release excludes benchmark images and pretrained weights rather than assigning them a new license.

    Guidelines:

    - The answer [N/A] means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.
    - The name of the license (e.g., CC-BY 4.0) should be included for each asset.
    - For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
    - If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
    - For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
    - If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [Yes]

    Justification: The accompanying code and numerical-artifact package provides a README, provenance and hash manifests, replay commands, and explicit descriptions of unavailable historical outputs. Appendices D and E identify the implemented protocols and evidence supporting the reported results.

    Guidelines:

    - The answer [N/A] means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: The research does not involve crowdsourcing, participant recruitment, or direct human-subject experiments; it analyzes synthetic data and existing benchmarks.

    Guidelines:

    - The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: The research does not recruit or interact with human participants and conducts no human-subject study requiring participant-risk disclosure or IRB approval.

    Guidelines:

    - The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
    - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
    - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an important, original, or non-standard component of the core research methods. The fixed CLIP nuisance direction described in Appendix C is a pretrained vision–language representation, not an LLM-based experimental method.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.